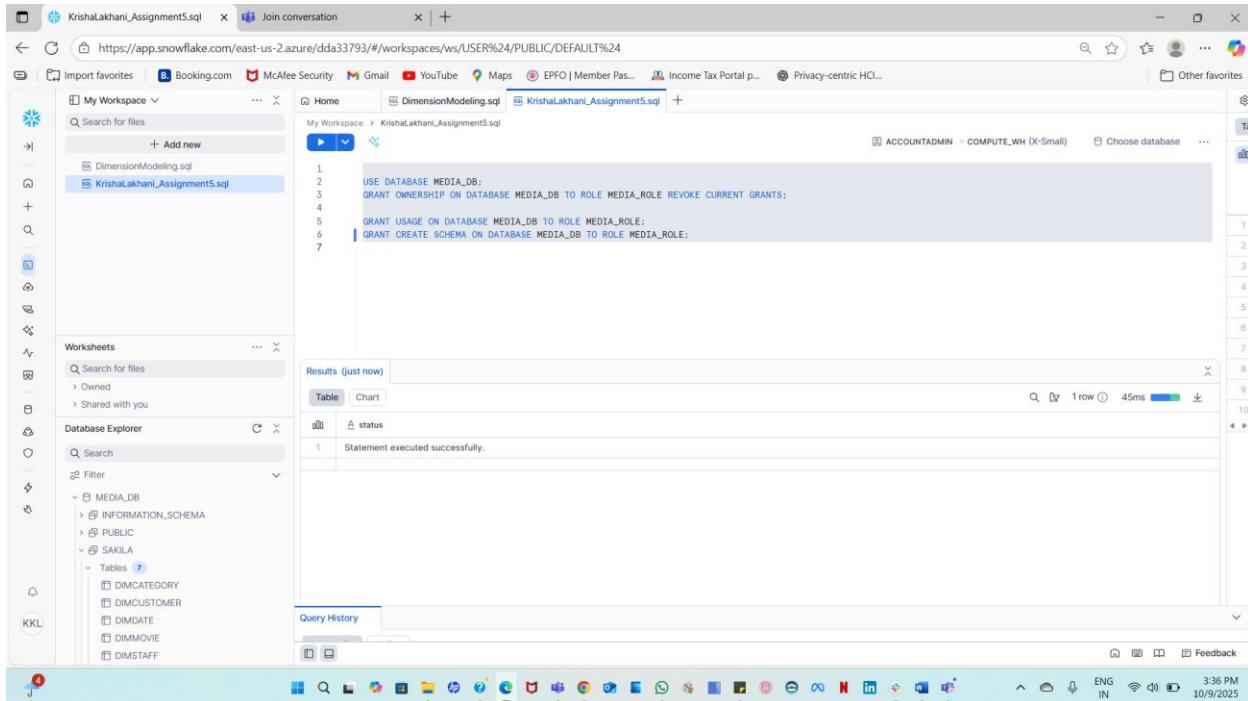


Assignment 5

Switch Role to ACCOUNTADMIN

Grant privileges to MEDIA_ROLE

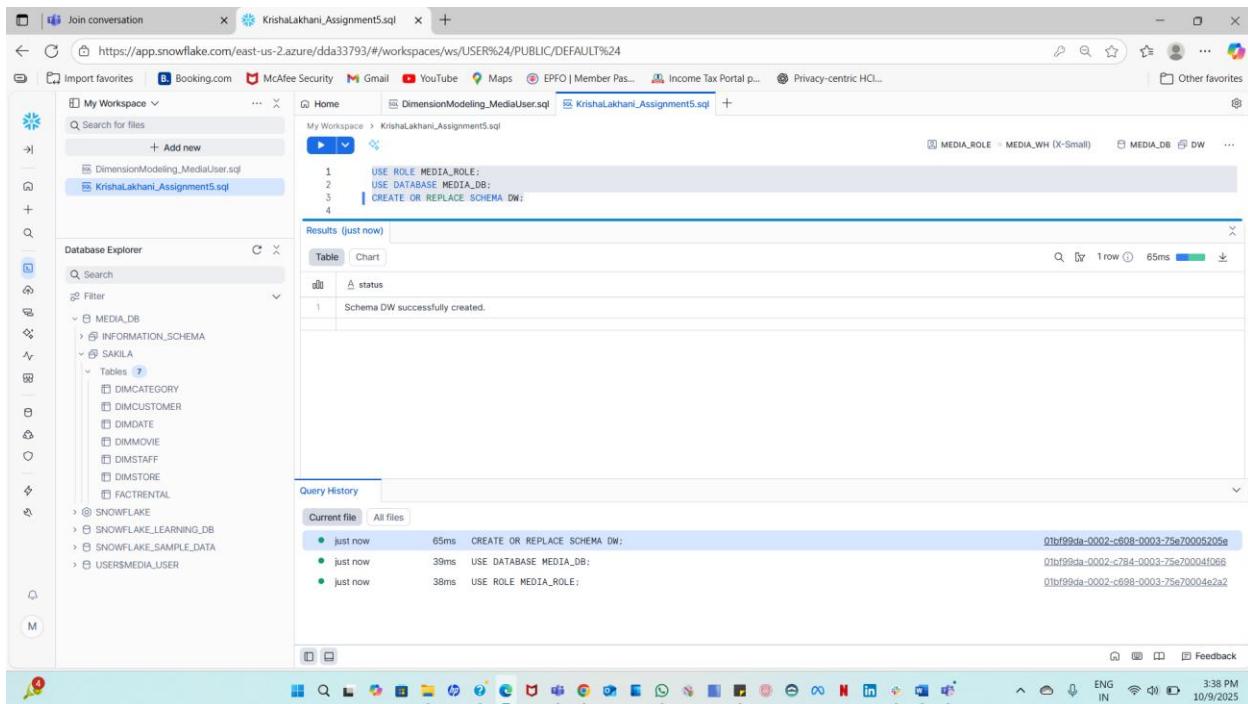


The screenshot shows the Snowflake web interface with a query editor window open. The URL is <https://app.snowflake.com/east-us-2.azure/dda33793/#/workspaces/ws/USER%24/PUBLIC/DEFAULT%24>. The query being run is:

```
1 USE DATABASE MEDIA_DB;
2 GRANT OWNERSHIP ON DATABASE MEDIA_DB TO ROLE MEDIA_ROLE REVOKE CURRENT GRANTS;
3
4 GRANT USAGE ON DATABASE MEDIA_DB TO ROLE MEDIA_ROLE;
5
6 GRANT CREATE SCHEMA ON DATABASE MEDIA_DB TO ROLE MEDIA_ROLE;
7
```

The results show one row: "Statement executed successfully."

Switch back to MEDIA_ROLE – Step 1



The screenshot shows the Snowflake web interface with a query editor window open. The URL is <https://app.snowflake.com/east-us-2.azure/dda33793/#/workspaces/ws/USER%24/PUBLIC/DEFAULT%24>. The query being run is:

```
1 USE ROLE MEDIA_ROLE;
2 USE DATABASE MEDIA_DB;
3 CREATE OR REPLACE SCHEMA DW;
```

The results show one row: "Schema DW successfully created."

The Query History pane shows three recent queries:

- just now 65ms CREATE OR REPLACE SCHEMA DW; 01bf99da-0002-c608-0003-75e70005205e
- just now 39ms USE DATABASE MEDIA_DB; 01bf99da-0002-c784-0003-75e70004f066
- just now 38ms USE ROLE MEDIA_ROLE; 01bf99da-0002-c698-0003-75e70004e2a2

Assignment 5

The screenshot shows a web browser window for the Snowflake interface. The URL is <https://app.snowflake.com/east-us-2.azure/dda33793/#/workspaces/ws/USER%24/PUBLIC/DEFAULT%24>. The left sidebar shows a workspace named "My Workspace" containing files like "DimensionModeling_MediaUser.sql" and "Krishalakhan_Assignment5.sql". The main area displays the following SQL code:

```
121 SOURCE_ID NUMBER(10),
122 DATE_TO_WAREHOUSE DATETIME DEFAULT CURRENT_TIMESTAMP()
123 );
124 --CLUSTER BY (SALE_DATE);
125
126
127
128 CREATE OR REPLACE SEQUENCE SALES_FACT_SEQ START = 1 INCREMENT = 1;
```

The results pane shows the output of the last query: "Sequence SALES_FACT_SEQ successfully created." The bottom status bar indicates the date as 10/9/2025 and the time as 3:39 PM.

The screenshot shows a web browser window for the Snowflake interface. The URL is <https://app.snowflake.com/east-us-2.azure/dda33793/#/workspaces/ws/USER%24/PUBLIC/DEFAULT%24>. The left sidebar shows a workspace named "My Workspace" containing files like "DimensionModeling_MediaUser.sql" and "Krishalakhan_Assignment5.sql". The main area displays the following SQL code:

```
129 CREATE OR REPLACE SEQUENCE SALES_FACT_SEQ START = 1 INCREMENT = 1;
130
131 -- validation
132 SHOW TABLES IN SCHEMA DW;
```

The results pane shows the output of the last query, listing tables in the DW schema:

name	database_name	schema_name	kind	comment	cluster_by	# rows	# bytes	owner
ARTIST_DIM	MEDIA_DB	DW	TABLE			0	0	MEDIA_ROLE
CUSTOMER_DIM	MEDIA_DB	DW	TABLE			0	0	MEDIA_ROLE
DATE_DIM	MEDIA_DB	DW	TABLE			0	0	MEDIA_ROLE
SALES_FACT	MEDIA_DB	DW	TABLE			0	0	MEDIA_ROLE
TIME_DIM	MEDIA_DB	DW	TABLE			0	0	MEDIA_ROLE

The bottom status bar indicates the date as 10/9/2025 and the time as 3:41 PM.

Assignment 5

The screenshot shows the Snowflake web interface. The left sidebar shows the 'Database Explorer' with the 'MEDIA_DB' database selected, containing tables like ARTIST_DIM, CUSTOMER_DIM, DATE_DIM, SALES_FACT, and TIME_DIM. The main area displays the SQL code for creating a sequence:

```
1 USE ROLE MEDIA_ROLE;
2 USE DATABASE MEDIA_DB;
3
4 CREATE OR REPLACE SCHEMA DW;
5 USE SCHEMA DW;
6
7 SELECT CURRENT_DATE(), CURRENT_TIMESTAMP();
8
9
10 CREATE OR REPLACE TABLE DW.DATE_DIM (
11     DATE_KEY NUMBER(10) PRIMARY KEY,
12     ...
13 );
```

The results pane shows the sequence was successfully created:

Table	Chart
status	
1	Sequence SALES_FACT_SEQ successfully created.

The bottom right corner shows the date and time: 4:02 PM 10/9/2025.

Step 2

The screenshot shows the Snowflake web interface. The left sidebar shows the 'Database Explorer' with the 'MEDIA_DB' database selected, containing tables like DIMCATEGORY, DIMCUSTOMER, DIMDATE, DIMMOVIE, DIMSTAFF, DIMSTORE, and FACTRENTAL. The main area displays the SQL code for creating a schema and tables:

```
134 --2
135 CREATE OR REPLACE SCHEMA STAGE;
136 CREATE OR REPLACE SCHEMA DW;
```

The results pane shows the schema was successfully created:

Table	Chart
status	
1	Schema STAGE successfully created.

The bottom right corner shows the date and time: 3:45 PM 10/9/2025.

Assignment 5

The screenshot shows the Snowflake web interface with the following details:

- Query Editor:** The current query is `Krishalakhani_Assignment5.sql`. The code shown is:246 TrackId INTEGER .
247 UnitPrice NUMBER(10, 2) .
248 Quantity INTEGER .
249 Created_By STRING(100) .
250 Created_Dt DATE
251);
252
253
254
255
256
- Results:** The results show a single row: "Table INVOICELINE successfully created."
- Query History:** The history shows five recent queries, all successful, with execution times ranging from 183ms to 210ms.
- Database Explorer:** The sidebar shows the database structure, including the `MEDIA_DB` schema which contains tables like `DIMCATEGORY`, `DIMCUSTOMER`, etc.

The screenshot shows the Snowflake web interface with the following details:

- Query Editor:** The current query is `Krishalakhani_Assignment5.sql`. The code shown is:258 --Validation
259 SHOW TABLES IN SCHEMA STAGE;
- Results:** The results show a table with 6 rows, listing tables in the `STAGE` schema:| | created_on | name | database_name | schema_name | kind | comment | cluster_by | # rows | # bytes | owner | refer |
| --- | --- | --- | --- | --- | --- | --- | --- | --- | --- | --- | --- |
| 1 | 2025-10-09 12:45:19.291 -0700 | ALBUM | MEDIA_DB | STAGE | TABLE | | | 0 | 0 | MEDIA_ROLE | 1 |
| 2 | 2025-10-09 12:45:16.897 -0700 | ARTIST | MEDIA_DB | STAGE | TABLE | | | 0 | 0 | MEDIA_ROLE | 1 |
| 3 | 2025-10-09 12:45:19.733 -0700 | CUSTOMER | MEDIA_DB | STAGE | TABLE | | | 0 | 0 | MEDIA_ROLE | 1 |
| 4 | 2025-10-09 12:45:17.372 -0700 | GENRE | MEDIA_DB | STAGE | TABLE | | | 0 | 0 | MEDIA_ROLE | 1 |
| 5 | 2025-10-09 12:45:20.814 -0700 | INVOICE | MEDIA_DB | STAGE | TABLE | | | 0 | 0 | MEDIA_ROLE | 1 |
| 6 | 2025-10-09 12:45:21.035 -0700 | INVOICELINE | MEDIA_DB | STAGE | TABLE | | | 0 | 0 | MEDIA_ROLE | 1 |
- Query History:** The history shows five recent queries, all successful, with execution times ranging from 183ms to 210ms.
- Database Explorer:** The sidebar shows the database structure, including the `MEDIA_DB` schema which contains tables like `DIMCATEGORY`, `DIMCUSTOMER`, etc.

Assignment 5

Step 3

3A. 1 - Create Storage account

The screenshot shows the Microsoft Azure portal interface for creating a new storage account. The URL is <https://portal.azure.com/?l=en-en&gb#%2fcreate%2fMicrosoft.StorageAccount-ARM>. The user is at the 'Create a storage account' step.

Project details:
Subscription: Azure for Students
Resource group: DamgFall2025

Instance details:
Storage account name: krishastorageadam
Region: (US) East US 2
Preferred storage type: Standard (radio button selected)
Performance: Standard (radio button selected)
Redundancy: Locally-redundant storage (LRS)

Security:
Require secure transfer for REST API operations: checked
Allow enabling anonymous access on individual containers: unchecked
Enable storage account key access: checked
Default to Microsoft Entra authorization in the Azure portal: unchecked
Minimum TLS version: Version 1.2
Permitted scope for copy operations (preview): From any storage account

Hierarchical Namespace:
Hierarchical namespace, complemented by Data Lake Storage Gen2 endpoint, enables file and directory semantics, accelerates big data analytics workloads, and enables access control lists (ACLs). [Learn more](#)

Access protocols:

Previous | Next | Review + create | Give feedback

5:24 PM 10/9/2025

Assignment 5

The screenshot shows two consecutive pages from the Microsoft Azure portal.

Page 1: Create a storage account

This page is titled "Create a storage account". It has tabs for Basics, Advanced, Networking, Data protection, Encryption, Tags, and Review + create. The Basics tab is selected. The configuration includes:

- Subscription: Azure for Students
- Resource group: DamgFall2025
- Location: East US 2
- Storage account name: krishastoragedamg
- Preferred storage type: Standard
- Performance: Standard
- Replication: Locally-redundant storage (LRS)

Page 2: Deployment Overview

This page shows the deployment status for "krishastoragedamg_1760045092027". The deployment is complete, indicated by a green checkmark icon. The deployment details are:

- Deployment name: krishastoragedamg_1760045092027
- Subscription: Azure for Students
- Resource group: DamgFall2025

The deployment started at 09/10/2025, 17:25:46. A correlation ID is provided: 0e4ba77-ade1-45fb-87e2-c66182edf878.

On the right side of the page, there are several promotional links:

- Cost Management**: Get notified to stay within your budget and prevent unexpected charges on your bill. [Set up cost alerts >](#)
- Microsoft Defender for Cloud**: Secure your apps and infrastructure. [Go to Microsoft Defender for Cloud >](#)
- Free Microsoft tutorials**: Start learning today. [Start learning today >](#)
- Work with an expert**: Azure experts are service provider partners who can help manage your assets on Azure and be your first line of support. [Find an Azure expert >](#)

Assignment 5

3A.2 – Create Key vault

A Create a key vault - Microsoft A x My Account x Replicate Workshop on Chinoi x +

portal.azure.com/?l=en-en-gb#create/Microsoft.KeyVault

Home > Key vaults > Create a key vault ...

Basics Access configuration Networking Tags Review + create

Azure Key Vault is a cloud service used to manage keys, secrets, and certificates. Key Vault eliminates the need for developers to store security information in their code. It allows you to centralize the storage of your application secrets which greatly reduces the chances that secrets may be leaked. Key Vault also allows you to securely store secrets and keys backed by Hardware Security Modules or HSMs. The HSMs used are Federal Information Processing Standards (FIPS) 140-2 Level 2 validated. In addition, key vault provides logs of all access and usage attempts of your secrets so you have a complete audit trail for compliance.

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription * Azure for Students

Resource group * Dangfall2025

Create new

Instance details

Key vault name * krisha-keyvault

Region * East US 2

Pricing tier * Standard

Recovery options

Soft delete protection will automatically be enabled on this key vault. This feature allows you to recover or permanently delete a key vault and secrets for the duration of the retention period. This protection applies to the key vault and the secrets stored within the key vault.

To enforce a mandatory retention period and prevent the permanent deletion of key vaults or secrets prior to the retention period elapsing, you can turn on purge protection. When purge protection is enabled, secrets cannot be purged by users or

Previous Next Review + create Give feedback

ENG IN 5:30 PM 10/9/2025

A Create a key vault - Microsoft A x My Account x Replicate Workshop on Chinoi x +

portal.azure.com/?l=en-en-gb#create/Microsoft.KeyVault

Home > Key vaults > Create a key vault ...

Access configuration Networking Tags Review + create

Configure data plane access for this key vault

To access a key vault in data plane, all callers (users or applications) must have proper authentication and authorization. Authentication establishes the identity of the caller. Authorization determines which operations the caller can execute. [Learn more](#)

Permission model

Grant data plane access by using a [Azure RBAC](#) or [Key Vault access policy](#)

Azure role-based access control (recommended)

Vault access policy

Resource access

Azure Virtual Machines for deployment

Azure Resource Manager for template deployment

Azure Disk Encryption for volume encryption

Access policies

Access policies enable you to have fine grained control over access to vault items. [Learn more](#)

Name	Email	Key Permissions	Secret Permissions	Certificate Permissions
Krishna Lakhani	lakhani.kri@northeastern.edu	Get, List, Update, Create, Import, Delete, Recover, Backup, Restor...	Get, List, Set, Delete, Recover, Backup, Restore	Get, List, Update, Create, Import, Delete, Recover, Backup, Restor...

Previous Next Review + create Give feedback

ENG IN 5:30 PM 10/9/2025

Assignment 5

The screenshot shows the Microsoft Azure portal interface. A new Key Vault named 'krisha-keyvault' is being created. The configuration includes:

- Subscription:** Azure for Students
- Resource group:** DamgFall2025
- Key vault name:** krisha-keyvault
- Region:** East US 2
- Pricing tier:** Standard
- Soft-delete:** Enabled
- Purge protection during retention period:** Disabled
- Days to retain deleted vaults:** 90 days

Access configuration:

- Azure Virtual Machines for deployment: Disabled
- Azure Resource Manager for template deployment: Disabled
- Azure Disk Encryption for volume encryption: Disabled
- Permission model: Vault access policy
- Access policies: 1

Networking:

- Connectivity method: Public endpoint (all networks)

At the bottom, there are 'Previous', 'Next', and 'Create' buttons.

The screenshot shows the deployment status of the 'krisha-keyvault' Key Vault. The deployment is complete, indicated by a green checkmark. Deployment details include:

- Deployment name: krisha-keyvault
- Subscription: Azure for Students
- Resource group: DamgFall2025

Deployment status: Start time : 09/10/2025, 17:33:41 Correlation ID : 0a44bd57-c7d1-4e17-b1ee-1b9ce3640153

Next steps: Go to resource, Give feedback, Tell us about your experience with deployment.

Right sidebar features:

- Cost management:** Get notified to stay within your budget and prevent unexpected charges on your bill. Set up cost alerts >
- Microsoft Defender for Cloud:** Secure your apps and infrastructure. Go to Microsoft Defender for Cloud >
- Free Microsoft tutorials:** Start learning today >
- Work with an expert:** Azure experts are service provider partners who can help manage your assets on Azure and be your first line of support. Find an Azure expert >

At the bottom, there is a message: Add or remove Favorites by pressing Ctrl + Shift F12.

Assignment 5

3A. 3 – Create ADF Service

The screenshot shows the 'Create Data Factory' wizard on the Microsoft Azure portal. The current step is 'Basics'. The page includes fields for Project details (Subscription: Azure for Students, Resource group: Dangfall2025), Instance details (Name: ADF-krisha-DW, Region: East US 2, Version: V2), and a 'Review + create' button at the bottom.

Project details

Select the subscription to manage deployed resources and costs. Use resource groups like folders to organize and manage all your resources.

Subscription: Azure for Students
Resource group: Dangfall2025
Create new

Instance details

Name: ADF-krisha-DW
Region: East US 2
Version: V2

Review + create

The screenshot shows the 'Create Data Factory' wizard on the Microsoft Azure portal. The current step is 'Git configuration'. It explains that Azure Data Factory allows you to configure a Git repository with either Azure DevOps or GitHub. It is a version control system that allows for easier change tracking and collaboration. There is a link to 'Learn more about Git integration in Azure Data Factory'. Below this, there is a checkbox for 'Configure Git later' which is checked.

Azure Data Factory allows you to configure a Git repository with either Azure DevOps or GitHub. Git is a version control system that allows for easier change tracking and collaboration.
Learn more about Git integration in Azure Data Factory

Configure Git later

Review + create

The screenshot shows the 'Create Data Factory' wizard on the Microsoft Azure portal. The current step is 'Review + create'. It displays the summary of the configuration: Name (ADF-krisha-DW), Region (East US 2), Version (V2), and Git configuration (Configure Git later checked). At the bottom, there are 'Previous', 'Next', and 'Review + create' buttons.

ADF-krisha-DW
East US 2
V2
Configure Git later

Review + create

Assignment 5

The screenshot shows the 'Create Data Factory' wizard on the 'Networking' tab. It includes sections for 'Managed virtual network' (checkbox for enabling managed virtual network), 'Self-hosted integration runtime inbound connectivity to Azure Data Factory service' (radio button for 'Public endpoint'), and a note about changing connectivity methods.

Managed virtual network
Choose whether you want the default AutoResolveIntegrationRuntime to be provisioned on demand inside an ADF-managed virtual network. If this setting is disabled, after the data factory is created, you can still choose whether to provision explicitly created Azure integration runtime inside an ADF-managed virtual network.
[Learn more](#)

Self-hosted integration runtime inbound connectivity to Azure Data Factory service
Choose whether to connect your self-hosted integration runtime to Azure Data Factory via public endpoint or private endpoint. This applies to self-hosted integration runtime running either on-premises or inside customer managed Azure virtual network.
[Learn more](#)

Connect via * Public endpoint Private endpoint

ⓘ You can change this or configure another connectivity method after this resource is created. [Learn more](#) ⓘ

Previous Next Review + create

The screenshot shows the 'Create Data Factory' wizard on the 'Review + create' tab. It displays the summary of the configuration and provides a 'Create' button to proceed.

Basics

Subscription	Azure for Students
Resource group	Damgfall2023
Name	ADF-Krishna-DW
Region	East US 2
Version	V2

Networking

Connect via Public endpoint

Previous Next Create

Assignment 5

The screenshot shows the Microsoft Azure portal with a deployment overview for 'Microsoft.DataFactory-20251009173651'. The deployment status is 'Your deployment is complete' with a green checkmark. Deployment details include a name of 'Microsoft.DataFactory-20251009173651', a subscription of 'Azure for Students', and a resource group of 'DamgFall2025'. The start time was 09/10/2025, 17:39:29, and the correlation ID is 741ed775-80db-4a85-86cc-6fbbe29e7656. On the right side, there are promotional cards for Cost management, Microsoft Defender for Cloud, Free Microsoft tutorials, and Work with an expert.

3B. Linked Services

Azure storage blob

The screenshot shows the 'Linked services' section in the Azure Data Factory blade. A new linked service is being created for 'Azure Blob Storage'. The 'Name' field is set to 'ls_azure_storage_blob'. The 'Connect via integration runtime' dropdown is set to 'AutoResolveIntegrationRuntime'. The 'Authentication type' is 'System-assigned managed identity'. The 'Account selection method' is 'From Azure subscription', and the selected subscription is 'Azure for Students (cbaeb871-8e5f-49d0-825b-6bcd75b282f1)'. The 'Storage account name' is 'krishastorageadmg' and the 'Storage account kind' is 'StorageV2'. At the bottom, there are 'Create' and 'Back' buttons, along with a 'Test connection' button.

Assignment 5

The screenshot shows the Microsoft Azure Data Factory interface. The left sidebar navigation includes General, Connections (selected), Source control, Author, and Security. Under Connections, Linked services is selected. The main content area displays a table titled "Linked services" with two entries:

Name	Type	Related	Annotations
ls_azure_key_vault	Azure Key Vault	0	
ls_azure_storage_blob	Azure Blob Storage	0	

Key vault

The screenshot shows the "New linked service" configuration page for an Azure Key Vault. The "Name" field is set to "ls_azure_key_vault". The "Azure key vault selection method" section has "From Azure subscription" selected. The "Azure subscription" dropdown is set to "Azure for Students (cbaeb871-8e5f-49d0-825b-6bcd75b282f1)". The "Azure key vault name" dropdown is set to "krisha-keyvault". The "Authentication method" dropdown is set to "System-assigned managed identity". A note at the bottom states: "Managed identity name: ADF-krisha-DW Click to copy Managed identity object ID: 03938f25-01e9-4672-ae1d-b8bf71709ab2 Grant Data Factory service managed identity access to your Azure Key Vault. Learn more". The "Create" button is visible at the bottom.

Assignment 5

The screenshot shows the Microsoft Azure Data Factory interface. The left sidebar navigation includes General, Connections (selected), Source control, Author, and Security. Under Connections, there are options for Linked services, Integration runtimes, Microsoft Purview, and ADF in Microsoft Fabric. The main content area displays the 'Linked services' section with a table listing one item: 'ls_azure_key_vault' of type 'Azure Key Vault'. A search bar at the top right says 'Search'.

SQL Server

The screenshot shows the Microsoft Azure Data Factory interface with the 'Edit linked service' dialog open. The dialog is titled 'Edit linked service' and shows the 'SQL Server' type selected. The 'Name' field is set to 'SqlServer1'. The 'Description' field is empty. The 'Connect via integration runtime' dropdown is set to 'IntegrationRuntime1'. The 'Version' dropdown is set to '2.0 (Recommended)'. The 'Server name' field contains 'LAPTOP-VUOFL100'. The 'Database name' field contains 'DAMG7370FALL2025'. The 'Authentication type' dropdown is set to 'SQL authentication'. The 'User name' field contains 'damg7370_fall25'. The 'Password' field is highlighted in blue, indicating it is the active field. The 'Always encrypted' checkbox is unchecked. At the bottom, there are 'Apply' and 'Cancel' buttons, and a 'Test connection' link.

Assignment 5

The screenshot shows the 'Linked services' configuration page in the Azure Data Factory interface. On the left, a sidebar lists various settings like General, Connections, and Source control. The main area displays a list of existing linked services:

Name	Type	Ref.
ls_azure_key_vault	Azure Key Vault	0
ls_azure_storage_blob	Azure Blob Storage	0
SqServer1	SQL Server	0

To the right, a detailed configuration pane for 'SqServer1' is open. It shows the connection type as 'SQL Server' and provides fields for 'User name' (diamg7370_fall25), 'Password' (Azure Key Vault), and 'Always encrypted'. Below this, there are sections for 'Additional connection properties', 'Annotations', and 'Parameters'. A table under 'Parameters' lists two entries: 'FullyQualifiedDomain' (String, LAPTOP-V0JOF100) and 'Username' (String, diamg7370_fall25). At the bottom, there are 'Apply' and 'Cancel' buttons, along with a 'Test connection' link.

This screenshot shows the 'Edit linked service' configuration for 'SqServer1'. The 'Name' field is set to 'SqServer1'. The 'Description' field is empty. Under 'Connect via integration runtime', 'IntegrationRuntime1' is selected. The 'Version' is set to '2.0 (Recommended)'. The 'Server name' is '@linkedService().FullyQualifiedDomainName'. The 'Database name' is 'DAMG7370FALL2025'. The 'Authentication type' is 'SQL authentication'. The 'User name' is '@linkedService().Username' and the 'Password' is 'Azure Key Vault'. The 'Always encrypted' checkbox is unchecked. A green checkmark indicates 'Connection successful'. At the bottom, there are 'Apply' and 'Cancel' buttons, along with a 'Test connection' link.

Assignment 5

Snowflake

The screenshot shows the Microsoft Azure Data Factory interface. On the left, a sidebar lists various settings like General, Connections, and Source control. The main area is titled 'Linked services' and shows three existing entries: 'ls_azure_key_vault' (Azure Key Vault), 'ls_azure_storage_blob' (Azure Blob Storage), and 'SqlServer1' (SQL Server). A 'New' button is visible. On the right, a detailed configuration pane is open for creating a new linked service to 'Snowflake V2'. The 'Name' field is set to 'Snowflake1'. The 'Type' dropdown is set to 'Snowflake'. Under 'Connect via integration runtime', 'Autodeskintegrationruntime' is selected. The 'Version' dropdown shows '1.1 (Recommended)'. The 'Account name' field contains '@linkedService().AccountName'. The 'Database' field contains '@linkedService().Database'. The 'Warehouse' field contains '@linkedService().Warehouse'. The 'Authentication type' is set to 'Basic'. The 'User name' field contains '@linkedService().Username'. The 'Password' dropdown is set to 'Azure Key Vault', and the password field contains '*****'. At the bottom, there are 'Create', 'Back', 'Test connection', and 'Cancel' buttons.

This screenshot shows the same Azure Data Factory interface after the linked service has been created. The 'New linked service' configuration pane is now closed, and the 'Snowflake1' entry is visible in the 'Linked services' list under the 'Type' column. The rest of the interface remains the same, with the sidebar and other linked services visible.

Assignment 5

The screenshot shows the Microsoft Azure Data Factory interface for the 'ADF-krisha-DW' dataset. The left sidebar contains navigation links for General, Connector upgrade advisor, Factory settings, Connections, Linked services, Integration runtimes, Microsoft Purview, ADF in Microsoft Fabric, Source control, Git configuration, ARM template, Author, Triggers, Global parameters, Data flow libraries, Security, Credentials, Customer managed key, Outbound rules, Managed private endpoint, and Workflow orchestration manager. The main area is titled 'Linked services' and displays four entries: 'ls_azure_key_vault' (Azure Key Vault), 'ls_azure_storage_blob' (Azure Blob Storage), 'Snowflake1' (Snowflake V2), and 'SqlServer1' (SQL Server). A search bar at the top right allows filtering by name and annotations.

Snowflake -stage (SAS)

The screenshot shows the Microsoft Azure Storage Container 'snowflake-stage' settings page. The left sidebar includes links for Overview, Diagnose and solve problems, Access Control (IAM), Settings, Shared access tokens (selected), Access policy, Properties, and Metadata. The main content area is titled 'Shared access tokens' and describes a Shared Access Signature (SAS) as a URI for restricted access to an Azure Storage container. It shows the 'Account key' signing method selected, a 'Key 1' signing key, and a 'None' stored access policy. Under 'Permissions', 'Read', 'Add', 'Create', 'Write', and 'List' are checked. 'Allowed IP addresses' and 'Allowed protocols' are also configured. A note at the bottom states: 'A shared access signature (SAS) is a URI that grants restricted access to an Azure Storage container. Use it when you want to grant access to storage account resources for a specific time range without sharing your storage account key.' A timestamp at the bottom right indicates the screenshot was taken at 12:40 PM on 10/11/2025.

Assignment 5

A shared access signature (SAS) is a URI that grants restricted access to an Azure Storage container. Use it when you want to grant access to storage account resources for a specific time range without sharing your storage account key. Learn more about creating an account SAS.

Account key User delegation key

Key 1

Stored access policy: None

Permissions: 5 selected

Start and expiry date/time:

- Start: 11/10/2025 12:32:26 (UTC-05:00) Eastern Time (US & Canada)
- Expiry: 18/10/2025 20:47:26 (UTC-05:00) Eastern Time (US & Canada)

Allowed IP addresses: for example, 168.1.5.65 or 168.1.5.65-168.1...

Allowed protocols: HTTPS only HTTPS and HTTP

Generate SAS token and URL

Add or remove Favorites by pressing **Ctrl+Shift+F2**

ENG IN 12:49 PM 10/11/2025

Blob SAS token: sp=racwl&st=2025-10-11T16:32:26Z&se=2025-10-19T00:47:26Z&spr=https&sv=2024-11-04&sr=c&sig=P4w7qSHuRD2aEk9JZMz976YRTEa4hQz0PlT9keAuAaE%3D

Blob SAS URL: https://krishastoragedamg.blob.core.windows.net/snowflake-stage?sp=racwl&st=2025-10-11T16:32:26Z&se=2025-10-19T00:47:26Z&spr=https&sv=2024-11-04&sr=c&sig=P4w7qSHuRD2aEk9JZMz976YRTEa4hQz0PlT9keAuAaE%3D

SAS Token - **sp=racwl&st=2025-10-11T16:32:26Z&se=2025-10-19T00:47:26Z&spr=https&sv=2024-11-04&sr=c&sig=P4w7qSHuRD2aEk9JZMz976YRTEa4hQz0PlT9keAuAaE%3D**

SAS URL - **https://krishastoragedamg.blob.core.windows.net/snowflake-stage?sp=racwl&st=2025-10-11T16:32:26Z&se=2025-10-19T00:47:26Z&spr=https&sv=2024-11-04&sr=c&sig=P4w7qSHuRD2aEk9JZMz976YRTEa4hQz0PlT9keAuAaE%3D**

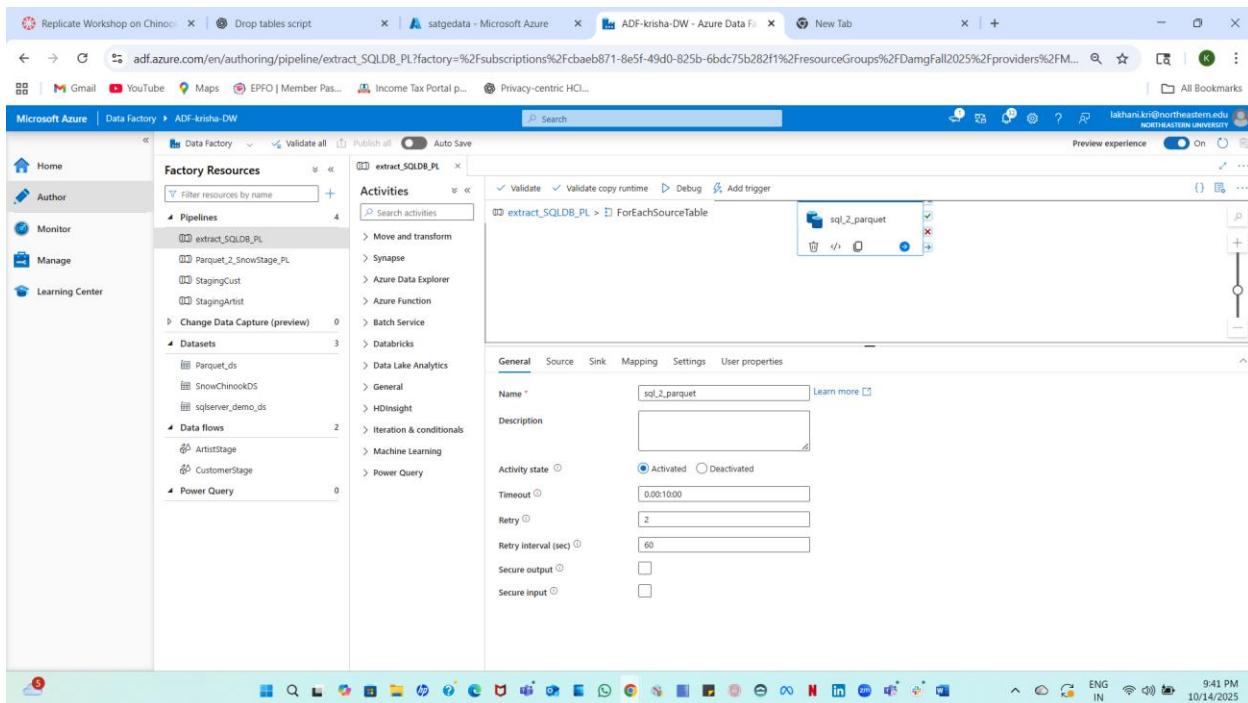
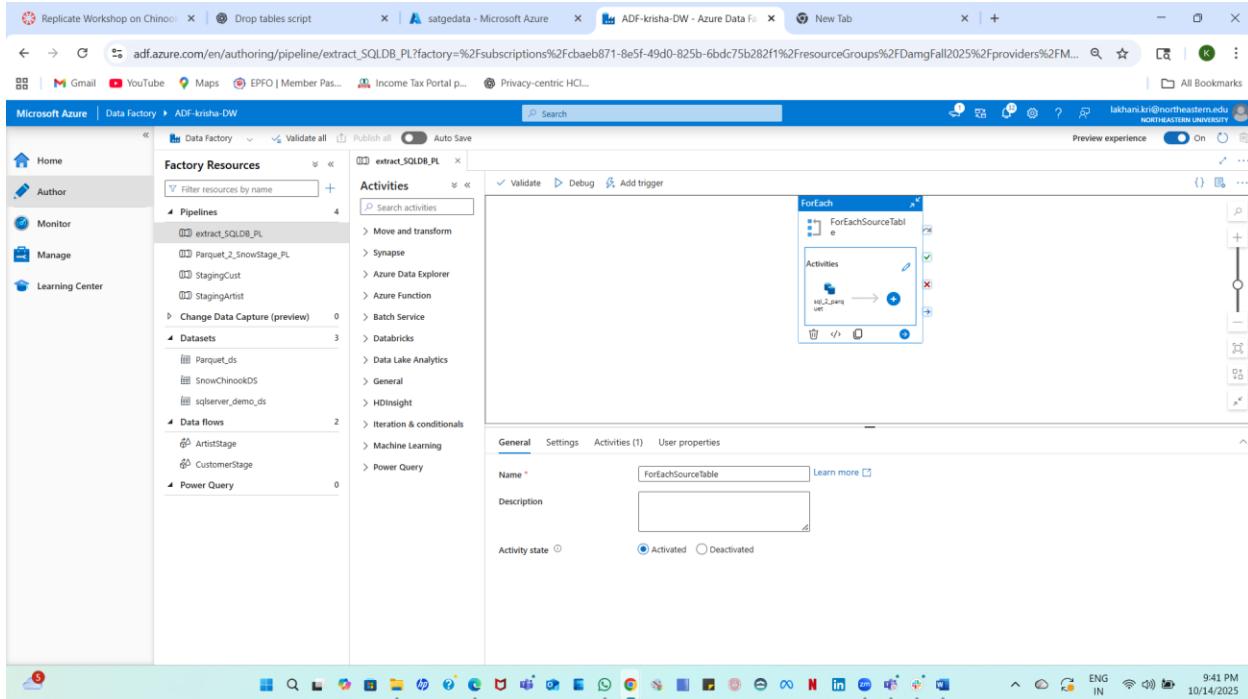
Assignment 5

The screenshot shows the Microsoft Azure Data Factory interface for the 'ADF-krisha-DW' dataset. On the left, the navigation menu includes options like General, Connections, and Linked services. Under 'Linked services', a list of existing connections is shown, including 'ls_azure_key_vault' (Azure Key Vault), 'ls_azure_storage_blob' (Azure Blob Storage), 'LS_Snowflake_SAS' (Snowflake V2), 'Snowflake1' (Snowflake V2), and 'SqlServer1' (SQL Server). A new connection is being configured for 'ls_azure_storage_blob'. The configuration pane shows the service type as 'Azure Blob Storage', authentication type as 'System-assigned managed identity', and service endpoint as 'https://krishasstorageadmg.blob.core.windows.net/'. The 'Test connection' button is visible at the bottom.

ADF

The screenshot shows the Microsoft Azure Data Factory pipeline editor for the 'ADF-krisha-DW' dataset. The left sidebar lists factory resources: Pipelines, Datasets, Data flows, and Power Query. A specific pipeline named 'extract_SQLDB_PL' is selected. The main workspace displays the pipeline's activities. One activity, 'ForEach', is expanded to show its sub-activities: 'ForEachSourceTable' and 'sql_Linq'. The pipeline parameters pane at the bottom shows a parameter named 'table_names_array' of type 'Array' with a default value of '["Invoice","InvoiceLine","Get"]'. The status bar indicates the date as 10/14/2025 and time as 9:41 PM.

Assignment 5



Assignment 5

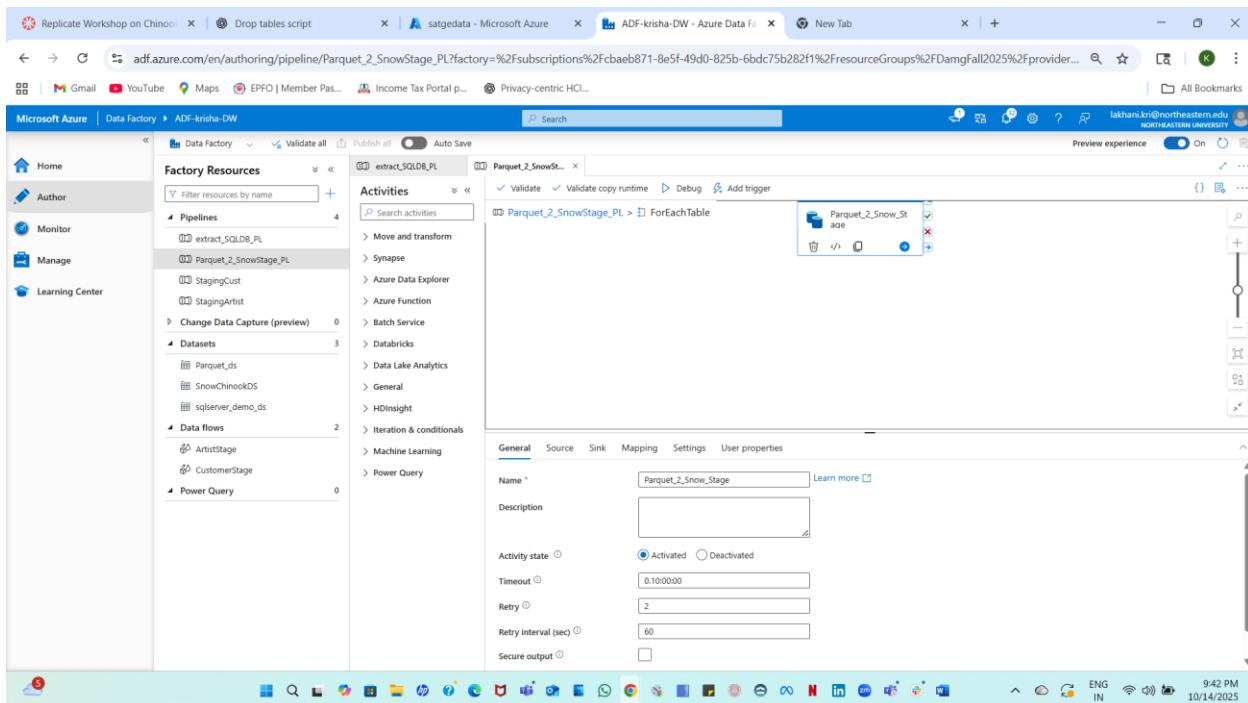
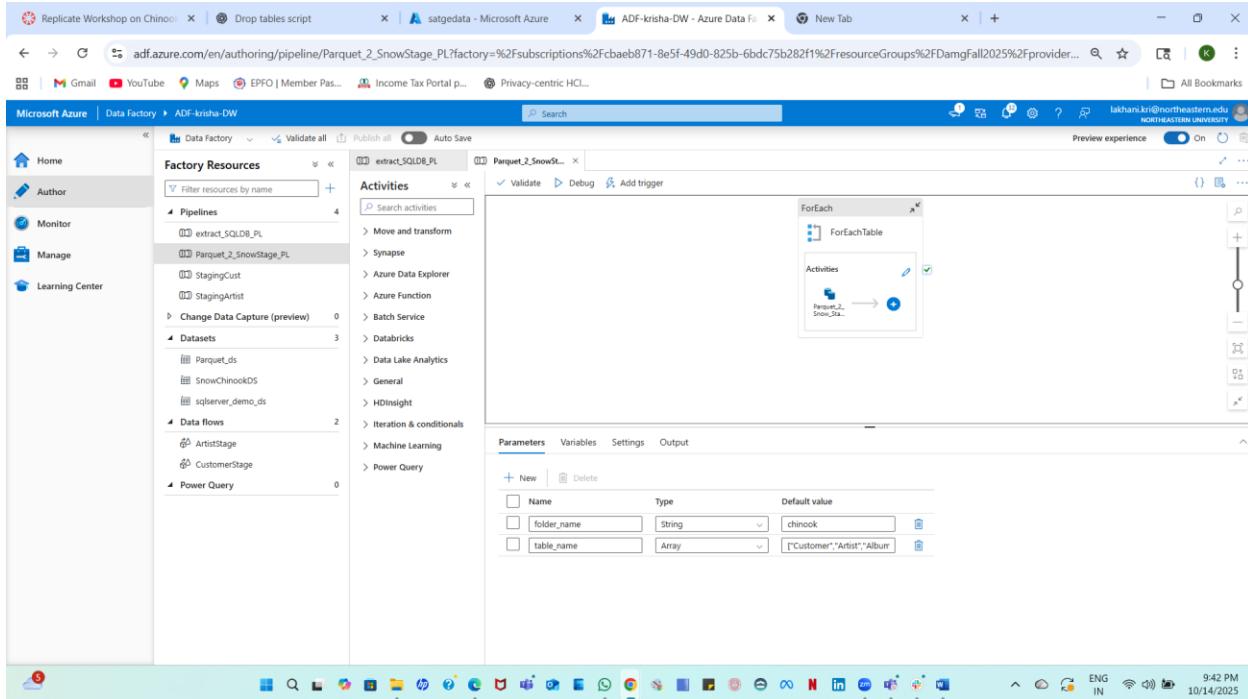
The screenshot shows the Azure Data Factory pipeline editor. On the left, the 'Factory Resources' sidebar lists Pipelines, Datasets, Data flows, and Power Query. The 'extract_SQLDB_Pl' pipeline is selected. In the main pane, the 'Activities' section shows a 'ForEachSourceTable' activity with a single child activity 'sql_2_parquet'. The 'Source' tab is active, displaying configuration for the source dataset:

- Source dataset:** sqserver_demo_d5
- Dataset properties:**
 - Name: schema_name, Value: chinook
 - Name: table_name, Value: @Item()
- Use query:** Table (radio button selected)
- Query timeout (minutes):** 120
- Isolation level:** Select...
- Partition option:** None (radio button selected)
- Additional columns:** + New

The screenshot shows the Azure Data Factory pipeline editor. The 'extract_SQLDB_Pl' pipeline is selected. In the main pane, the 'Activities' section shows a 'ForEachSourceTable' activity with a single child activity 'sql_2_parquet'. The 'Sink' tab is active, displaying configuration for the sink dataset:

- Sink dataset:** Parquet_d5
- Dataset properties:**
 - Name: container_name, Value: satgedata
 - Name: folder_name, Value: chinook
 - Name: file_name, Value: @concat(item(),'parquet')
- Copy behavior:** Select...
- Max concurrent connections:** (empty input field)
- Block size (MB):** (empty input field)
- Metadata:** + New
- Max rows per file:** (empty input field)

Assignment 5



Assignment 5

The screenshot shows the Azure Data Factory pipeline configuration for the 'extract_SQLDB_PL' pipeline. The pipeline consists of two activities: 'Parquet_2_SnowStage_PL' and 'Parquet_2_SnowStage_PL'. The first activity is a 'ForEachTable' activity that reads from a 'Parquet_ds' dataset. The second activity is a 'ForEachTable' activity that writes to a 'SnowChinookDS' dataset. Both activities have their 'Validate copy runtime' option checked.

Activities:

- Parquet_2_SnowStage_PL (For Each Table)
- Parquet_2_SnowStage_PL (For Each Table)

General Properties:

- Source dataset:** Parquet_ds
- Dataset properties:**
 - Name: container_name, Value: satgatedata
 - Name: folder_name, Value: @pipeline().parameters.folder_name
 - Name: file_name, Value: @concat(item().parquet*)
- File path type:** File path in dataset
- Filter by last modified:** Start time (UTC) [] End time (UTC) []
- Recursively:** Checked
- Enable partitions discovery:** Unchecked
- Max concurrent connections:** []
- Additional columns:** []

The screenshot shows the Azure Data Factory pipeline configuration for the 'extract_SQLDB_PL' pipeline, specifically focusing on the 'Sink' tab for the second 'Parquet_2_SnowStage_PL' activity. The sink dataset is 'SnowChinookDS'. The 'Sink properties' section includes 'schema_name: STAGE' and 'table_name: @toUpper(item())'. The 'Additional Snowflake copy options' section contains a single entry: 'ON_ERROR: CONTINUE'.

General Properties:

- Sink dataset:** SnowChinookDS
- Dataset properties:**
 - Name: schema_name, Value: STAGE
 - Name: table_name, Value: @toUpper(item())
- Pre-copy script:** []
- Storage integration:** Select or type... [] Refresh []
- Additional Snowflake copy options:**

Property name	Value
ON_ERROR	CONTINUE

Assignment 5

The screenshot shows the Microsoft Azure Storage Container Overview page for the 'satgedata' container. The 'chinook' blob is listed with its details: Last modified (10/13/2025, 11:35:50 AM), Access tier (Hot (inferred)), Blob type (Block blob), Size (9.4 kB), and Lease state (Available). The blob's URL is also provided.



The screenshot shows the Microsoft Azure Storage Container Overview page for the 'satgedata' container, specifically for the 'chinook' folder. The blobs listed are: i-1, Album.parquet, Artist.parquet, Customer.parquet, Genre.parquet, Invoice.parquet, and Invoiceline.parquet. Each blob has its last modified date, access tier, blob type, size, and lease state.

Name	Last modified	Access tier	Blob type	Size	Lease state
i-1	10/13/2025, 11:35:50 AM	Hot (inferred)	Block blob	9.4 kB	Available
Album.parquet	10/13/2025, 11:35:18 AM	Hot (inferred)	Block blob	6.49 kB	Available
Artist.parquet	10/13/2025, 11:35:18 AM	Hot (inferred)	Block blob	8.9 kB	Available
Customer.parquet	10/13/2025, 11:35:18 AM	Hot (inferred)	Block blob	892 B	Available
Genre.parquet	10/14/2025, 11:38:16 PM	Hot (inferred)	Block blob	13.98 kB	Available
Invoice.parquet	10/14/2025, 11:37:54 PM	Hot (inferred)	Block blob	22.08 kB	Available
Invoiceline.parquet	10/14/2025, 11:37:54 PM	Hot (inferred)	Block blob	13.98 kB	Available



Assignment 5

The screenshot shows the Microsoft Azure Data Factory interface for the 'ADF-krisha-DW' dataset. The left sidebar lists resources: Pipelines, Datasets, Data flows, and Power Query. The 'Parquet_ds' dataset is selected. The main pane shows the 'Parquet' icon and the dataset details. Under 'Connection', it is linked to 'is_azure_storage_blob'. The 'File path' is set to '@dataset().container_name' / '@dataset().folder_name' / '@dataset().file_name'. The 'Compression type' is set to 'snappy'. The status bar at the bottom indicates '9:45 PM 10/14/2025'.

This screenshot is identical to the one above, but the 'Parameters' tab is selected in the top navigation bar. It shows three parameters: 'container_name' (String, Value), 'folder_name' (String, Value), and 'file_name' (String, Value). The status bar at the bottom indicates '9:45 PM 10/14/2025'.

Assignment 5

The screenshot shows the Microsoft Azure Data Factory interface. On the left, the 'Factory Resources' sidebar lists Pipelines, Datasets, Data flows, and Power Query. The 'Datasets' section is expanded, showing 'Parquet_ds' and 'SnowChinookDS'. The 'SnowChinookDS' dataset is selected. In the main pane, the 'Parameters' tab is active. It shows a 'Linked service properties' table with the following data:

Name	Type
AccountName	string
Database	string
Warehouse	string
Username	string
Role	string
Host	string

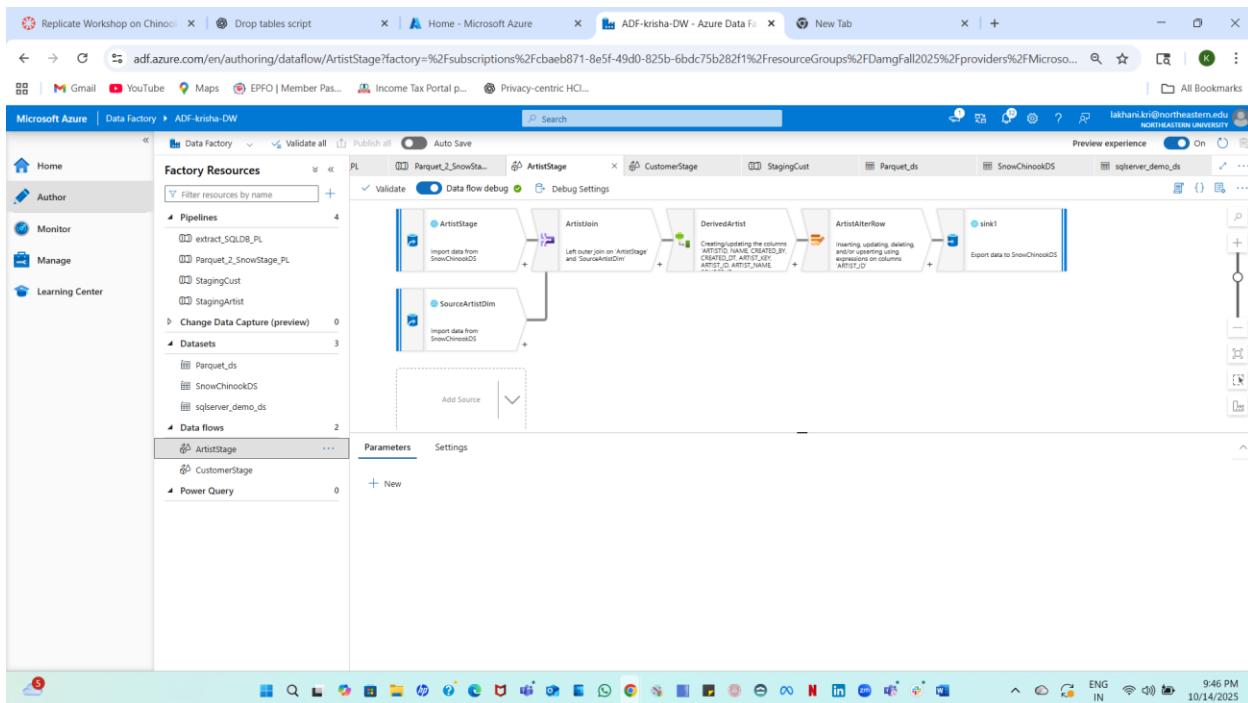
Below the table, there are input fields for 'schema_name' (@dataset().schema_name) and 'table_name' (@dataset().table_name), both with a 'Preview data' button. A checkbox for 'Enter manually' is also present.

This screenshot shows the same Azure Data Factory interface as the previous one, but with additional parameters added to the 'Parameters' tab. The 'Parameters' tab now contains two entries:

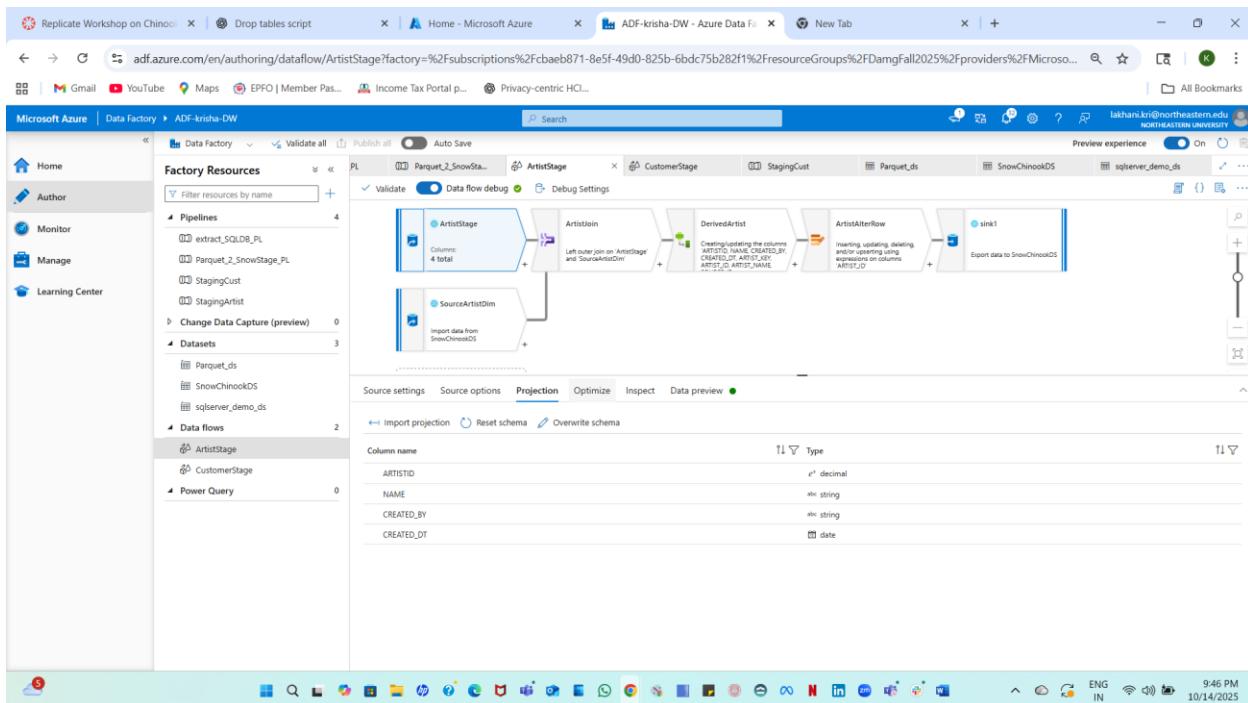
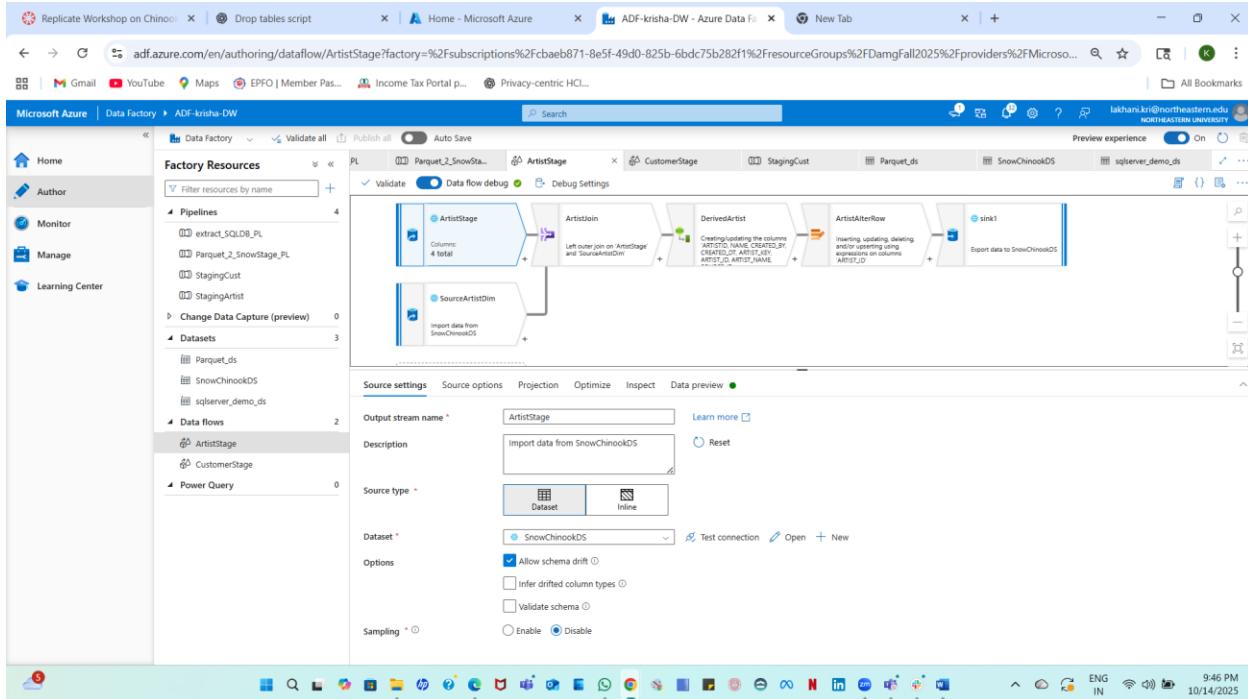
Name	Type	Default value
schema_name	String	STAGE
table_name	String	CUSTOMER

Assignment 5

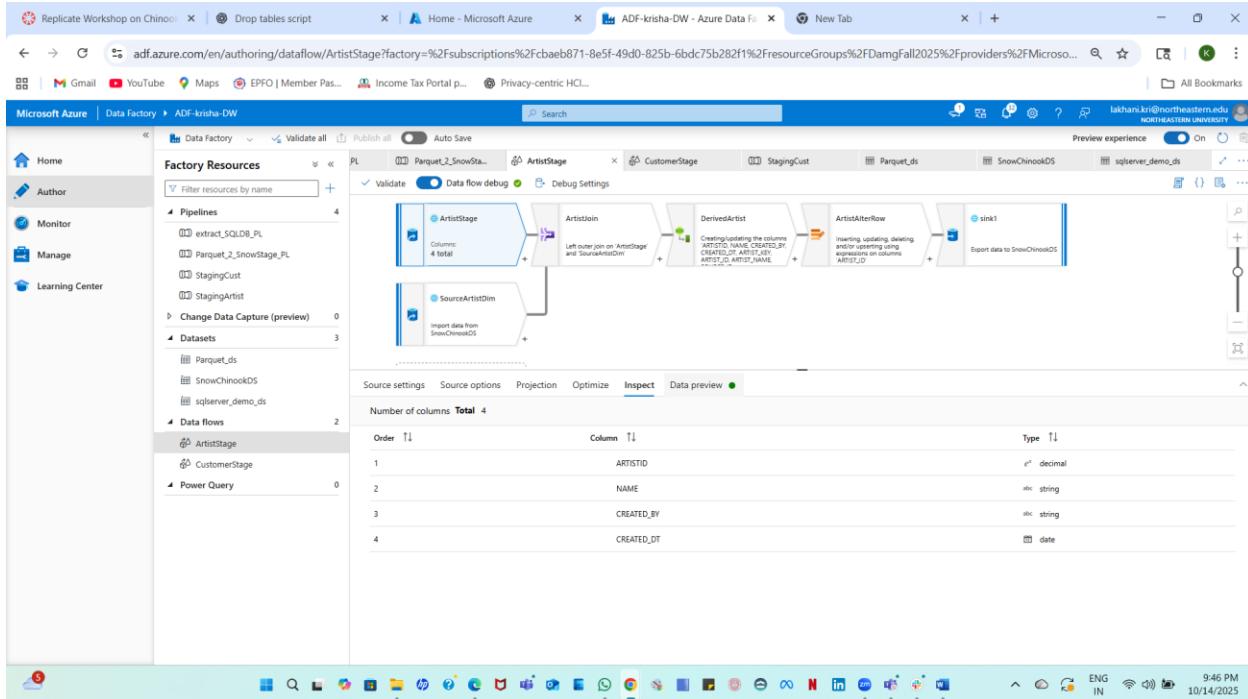
The screenshot shows the Microsoft Azure Data Factory interface for the 'ADF-krisha-DW' dataset. On the left, the 'Factory Resources' sidebar lists Pipelines, Datasets, and Data flows. The 'sqserver_demo.ds' dataset is selected. In the main pane, the 'Connection' tab is active, showing a linked service named 'SqlServer1' connected to a SQL Server database 'sqserver_demo.ds'. The 'Integration runtime' dropdown is set to 'integrationruntime1'. The 'Table' section displays parameters: '@dataset().schema_name' and '@dataset().table_name', both with a checked 'Enter manually' option.



Assignment 5



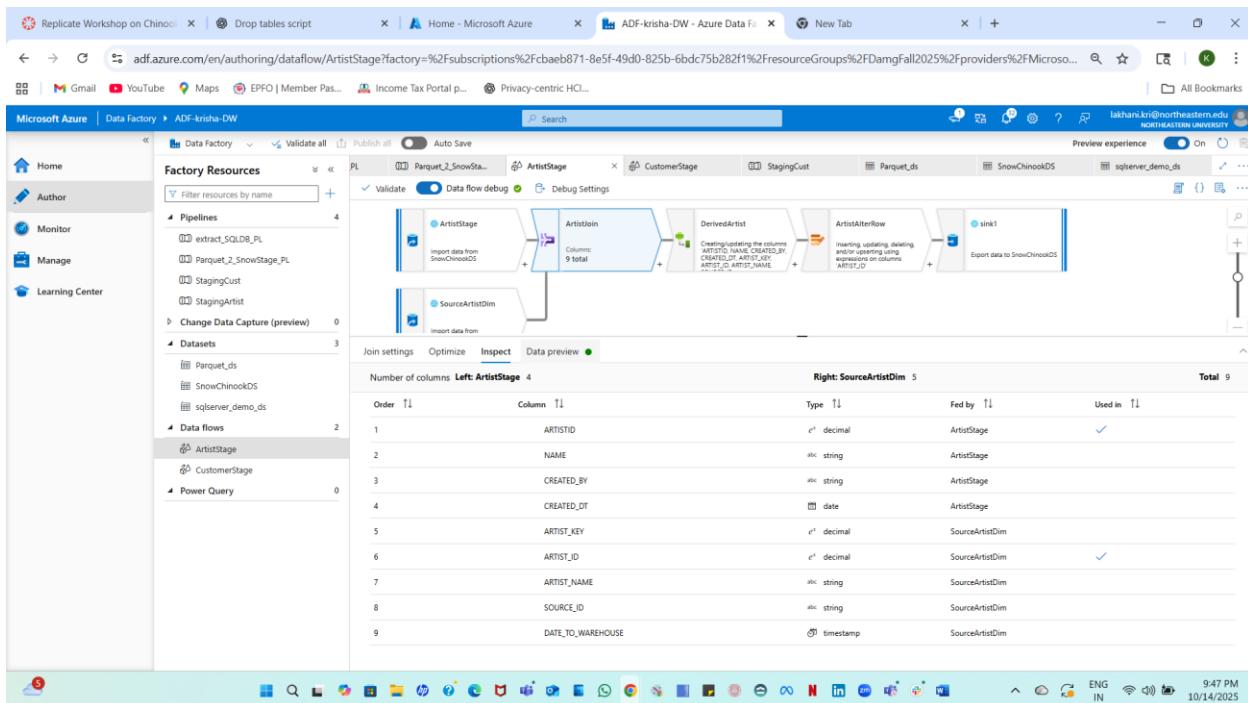
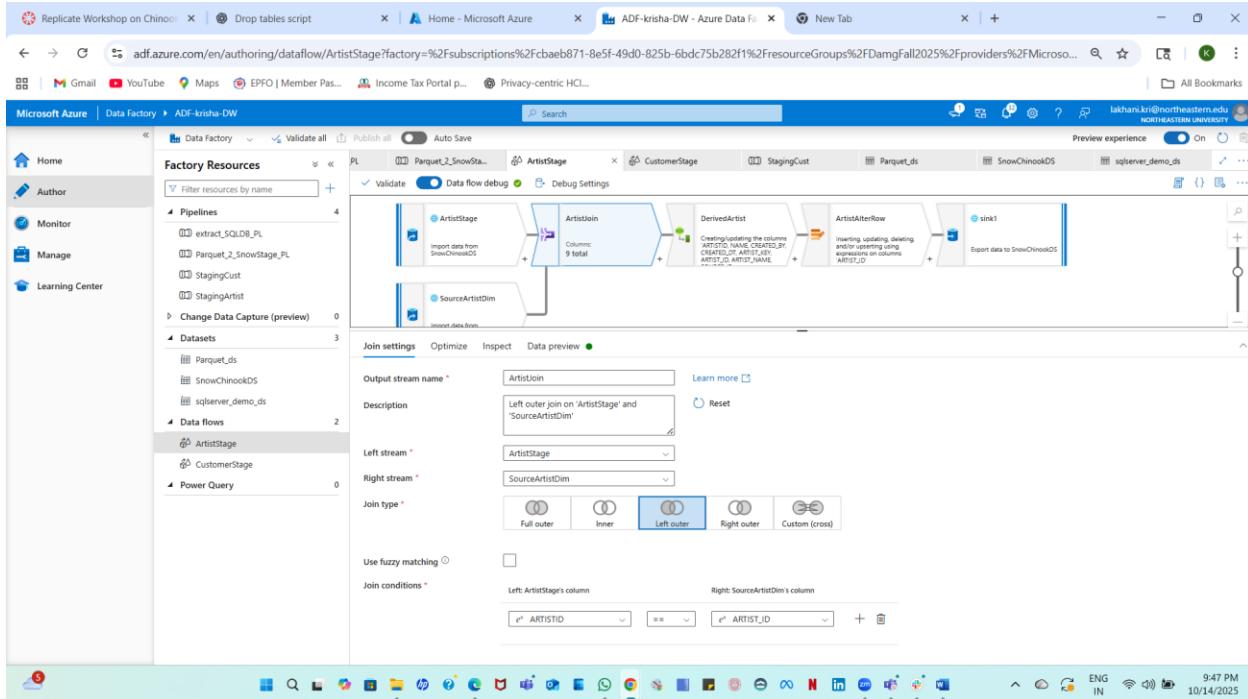
Assignment 5



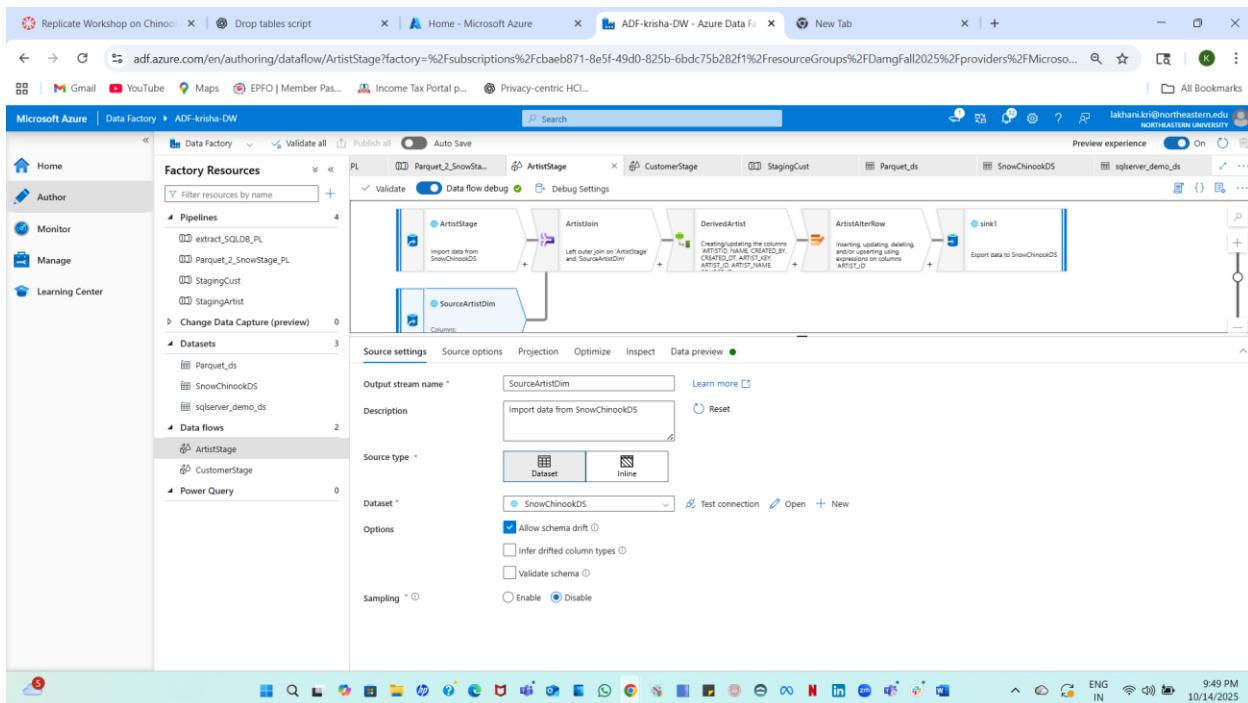
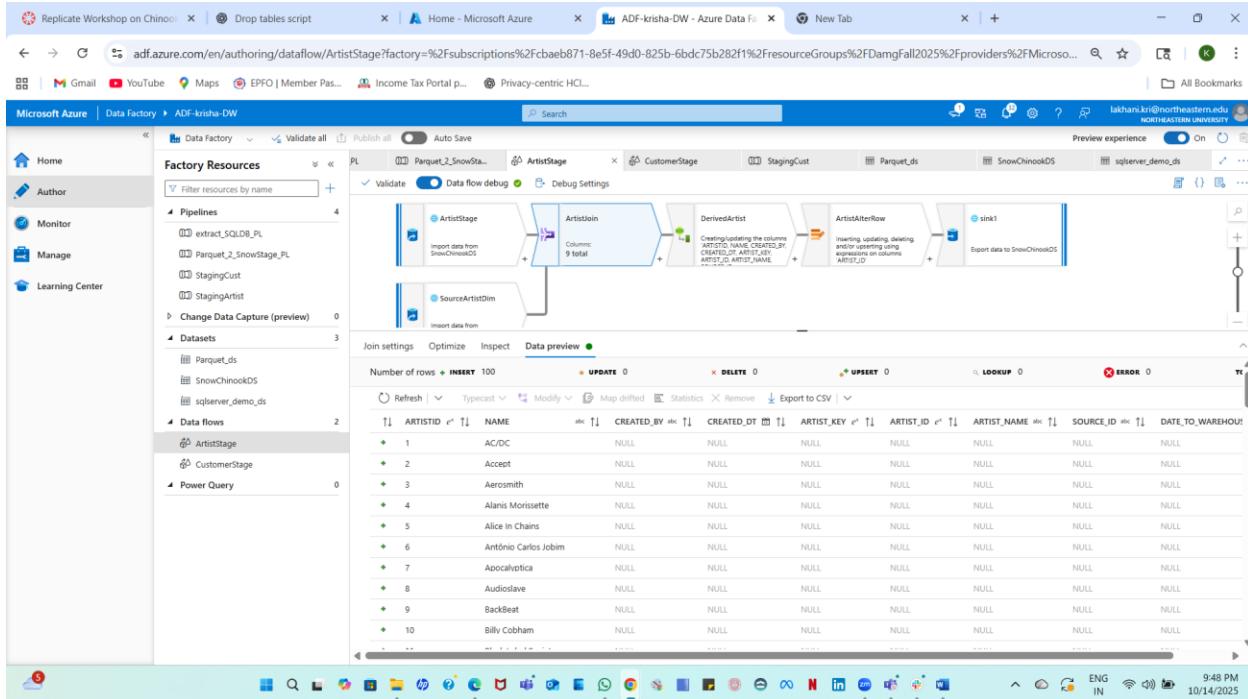
The screenshot shows the Microsoft Azure Data Factory Data Flow interface with the 'Data preview' tab selected. The preview table shows the results of the data flow, with 100 rows inserted. The columns are ARTISTID, NAME, CREATED_BY, and CREATED_DT. The data includes various artists like AC/DC, Accept, Aerosmith, etc., with null values for some fields.

ARTISTID	NAME	CREATED_BY	CREATED_DT
1	AC/DC	NULL	NULL
2	Accept	NULL	NULL
3	Aerosmith	NULL	NULL
4	Alanis Morissette	NULL	NULL
5	Alice In Chains	NULL	NULL
6	Antônio Carlos Jobim	NULL	NULL
7	Apocalyptica	NULL	NULL
8	Audioslave	NULL	NULL
9	BackBeat	NULL	NULL

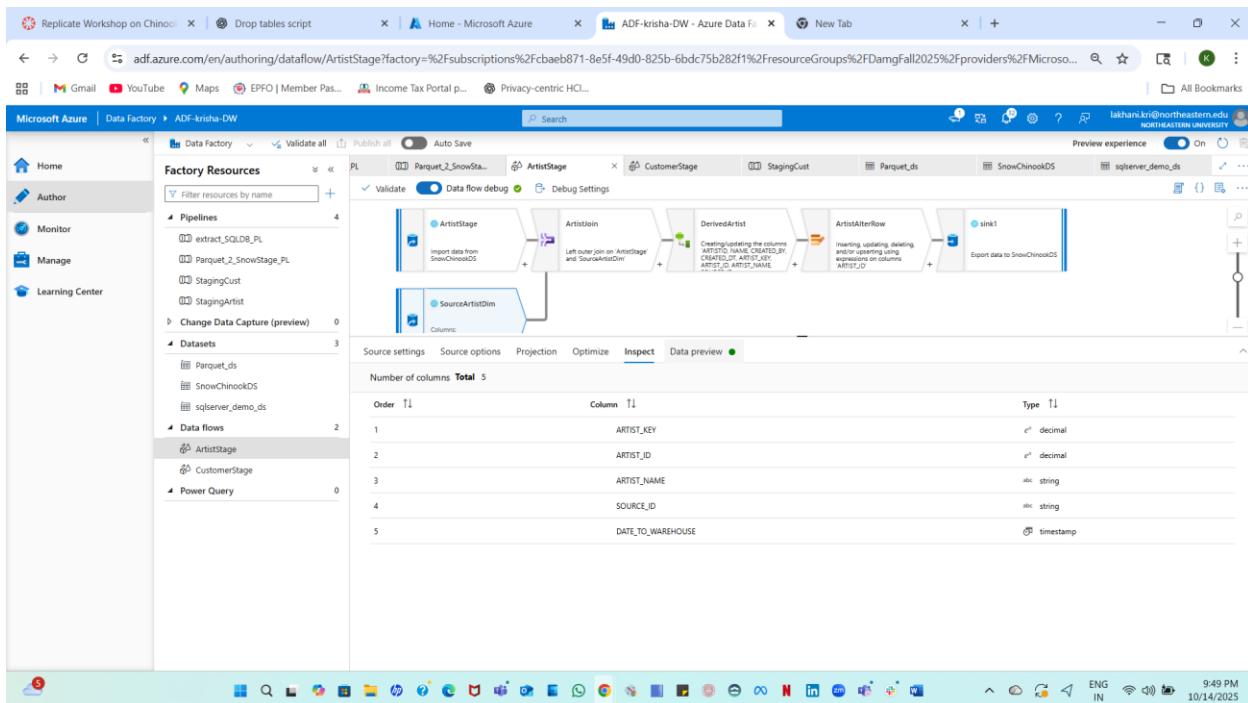
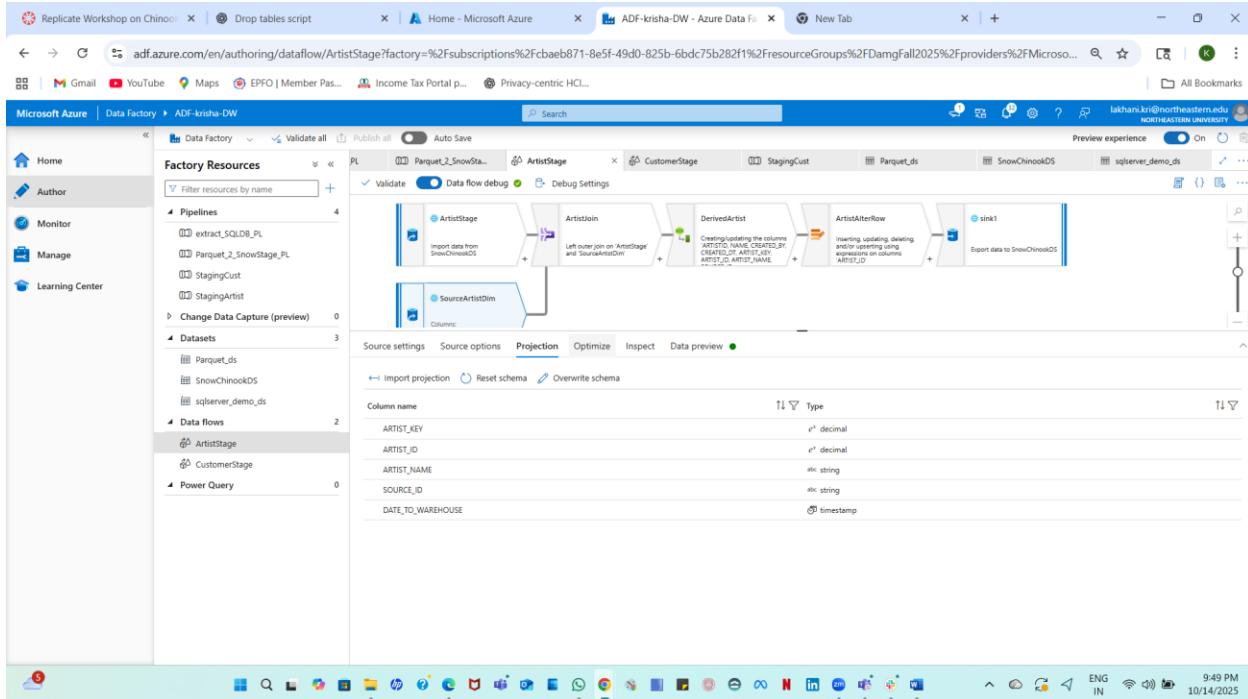
Assignment 5



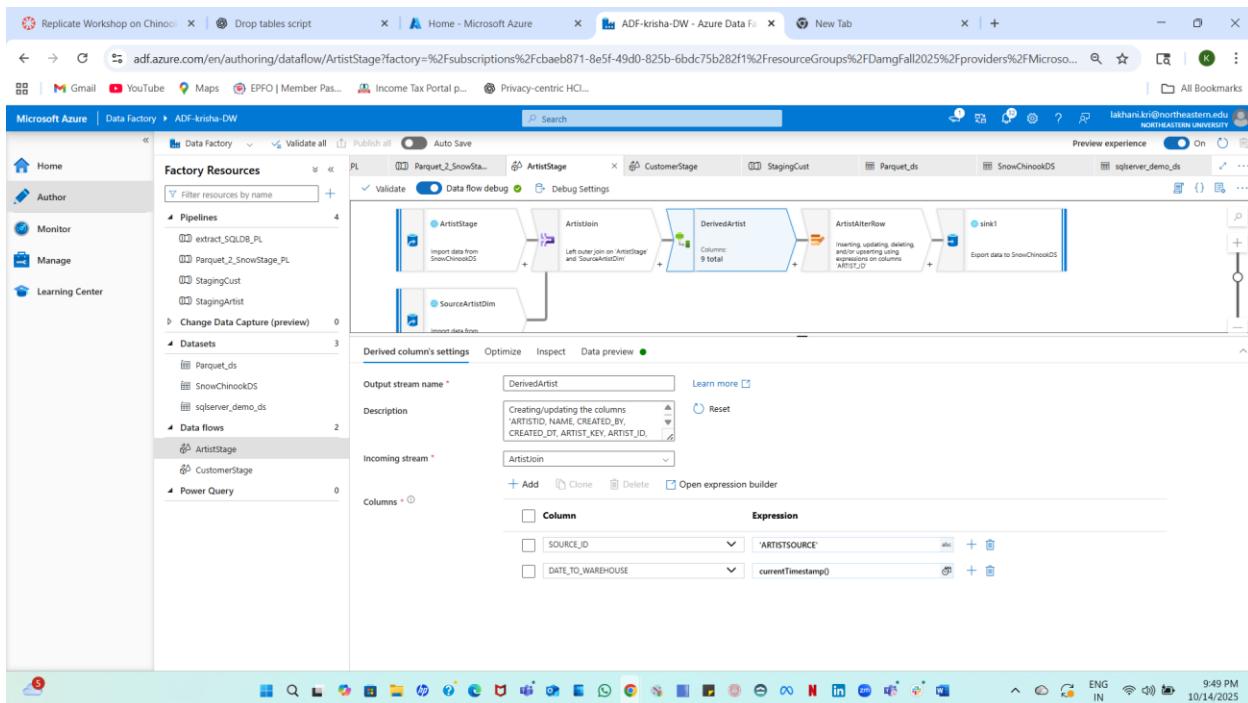
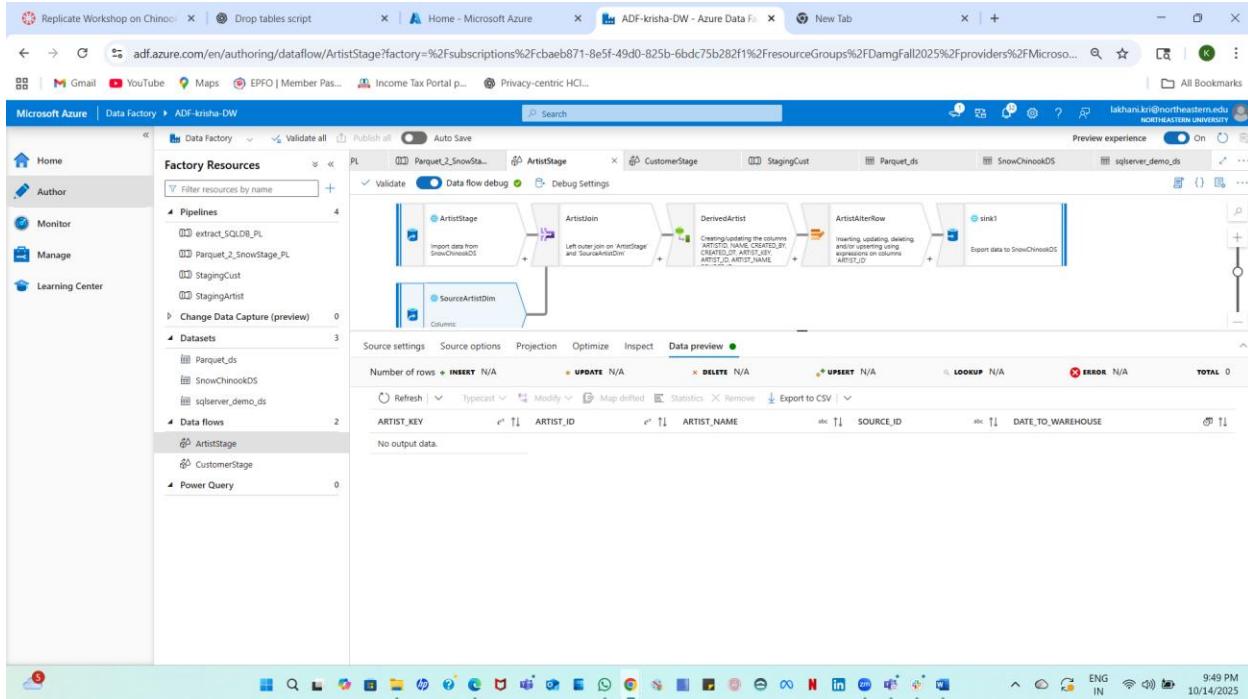
Assignment 5



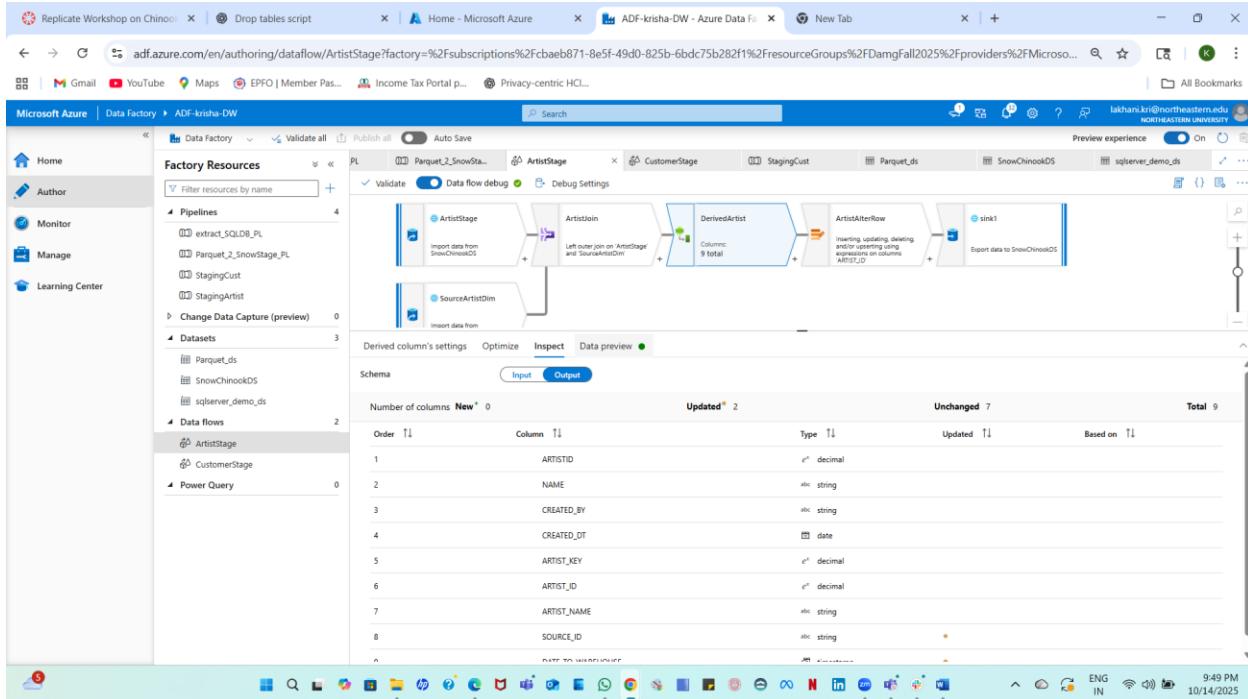
Assignment 5



Assignment 5



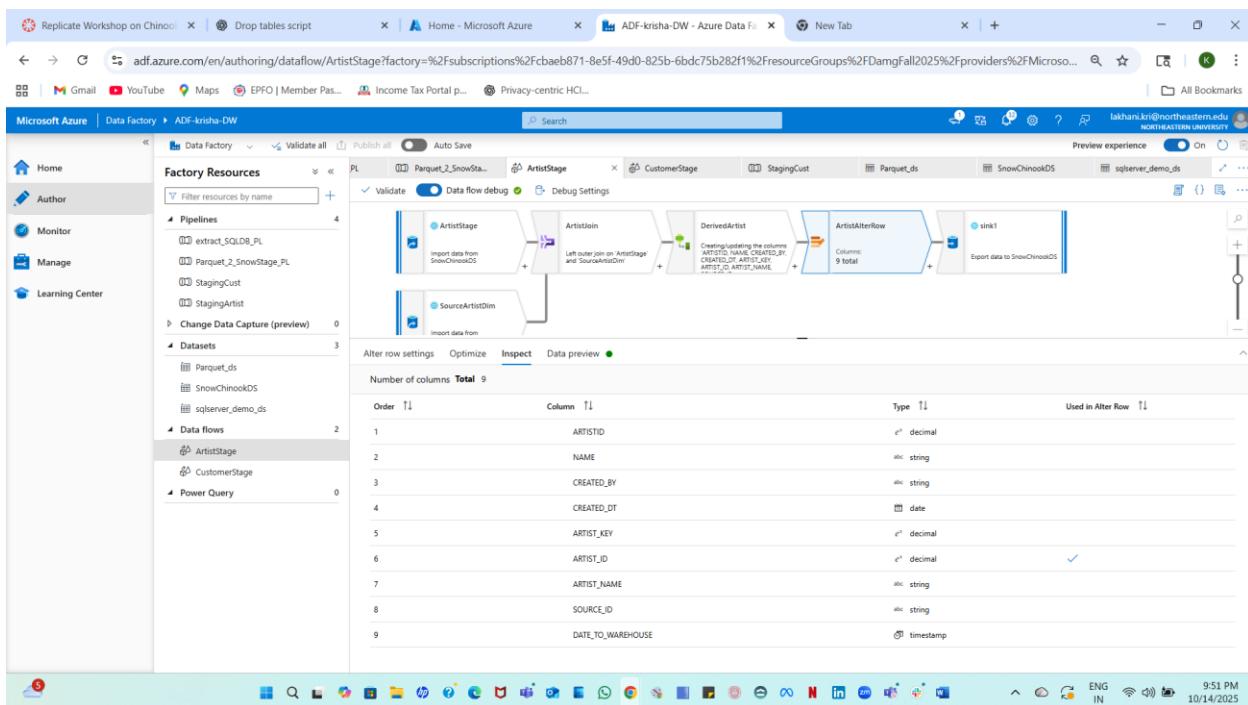
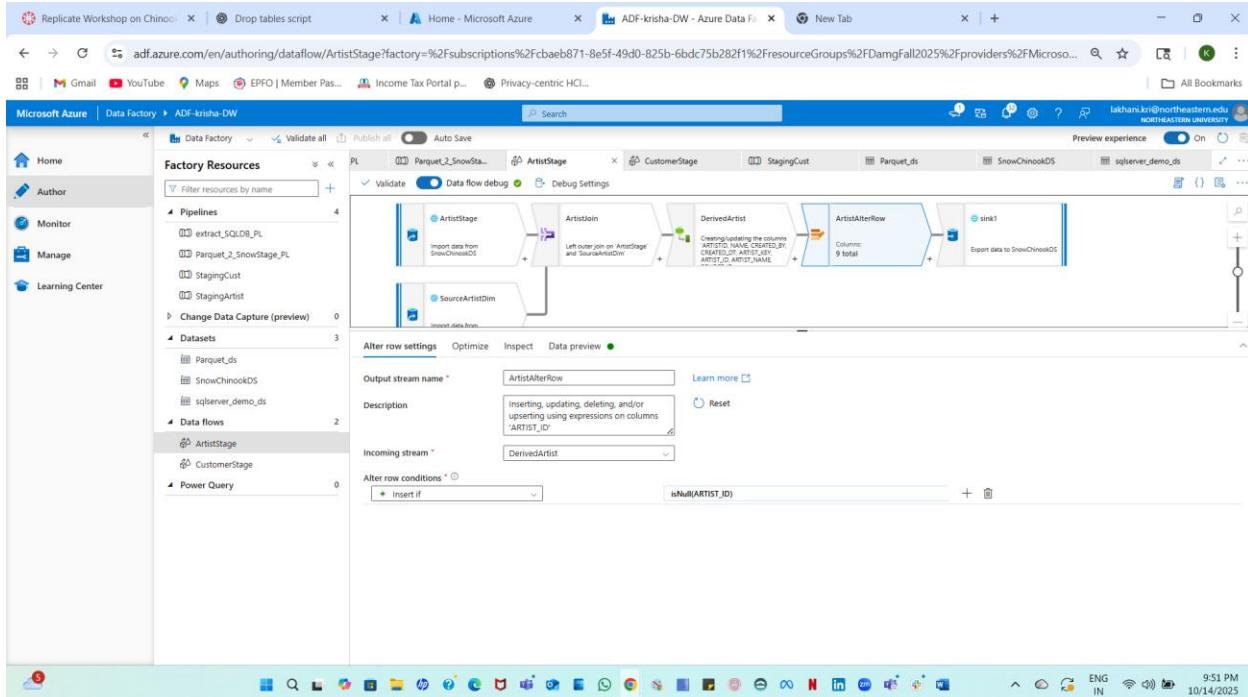
Assignment 5



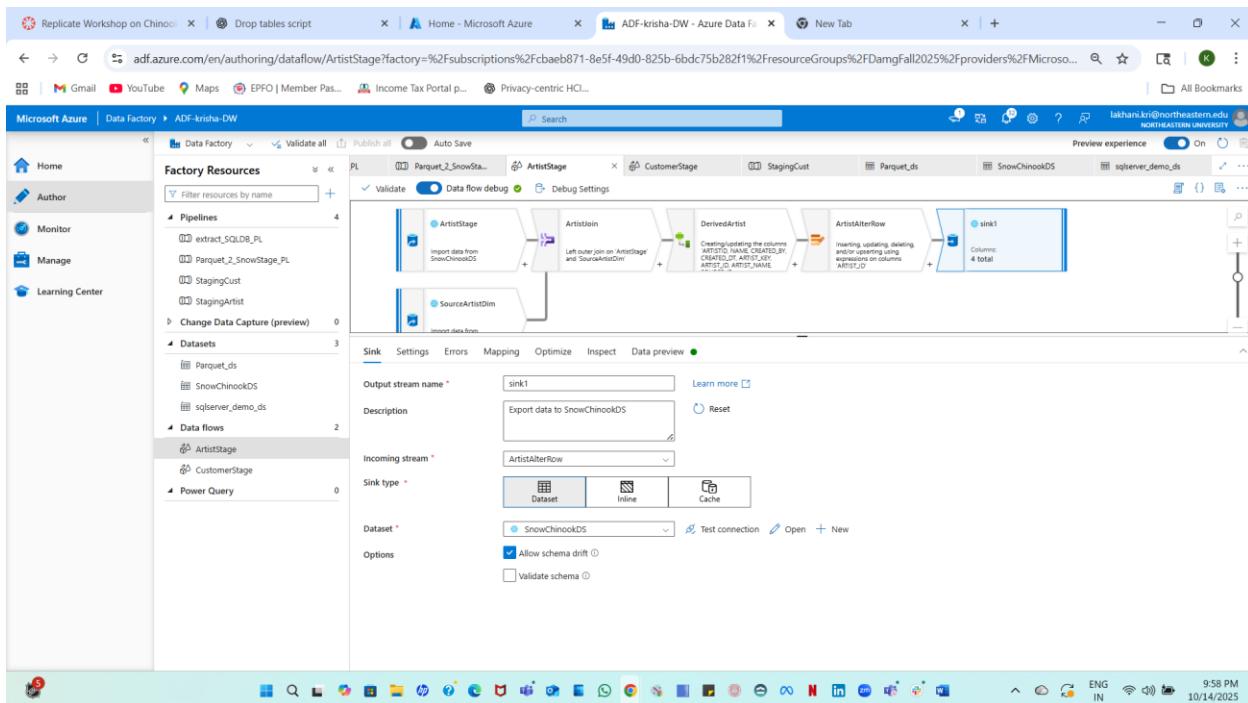
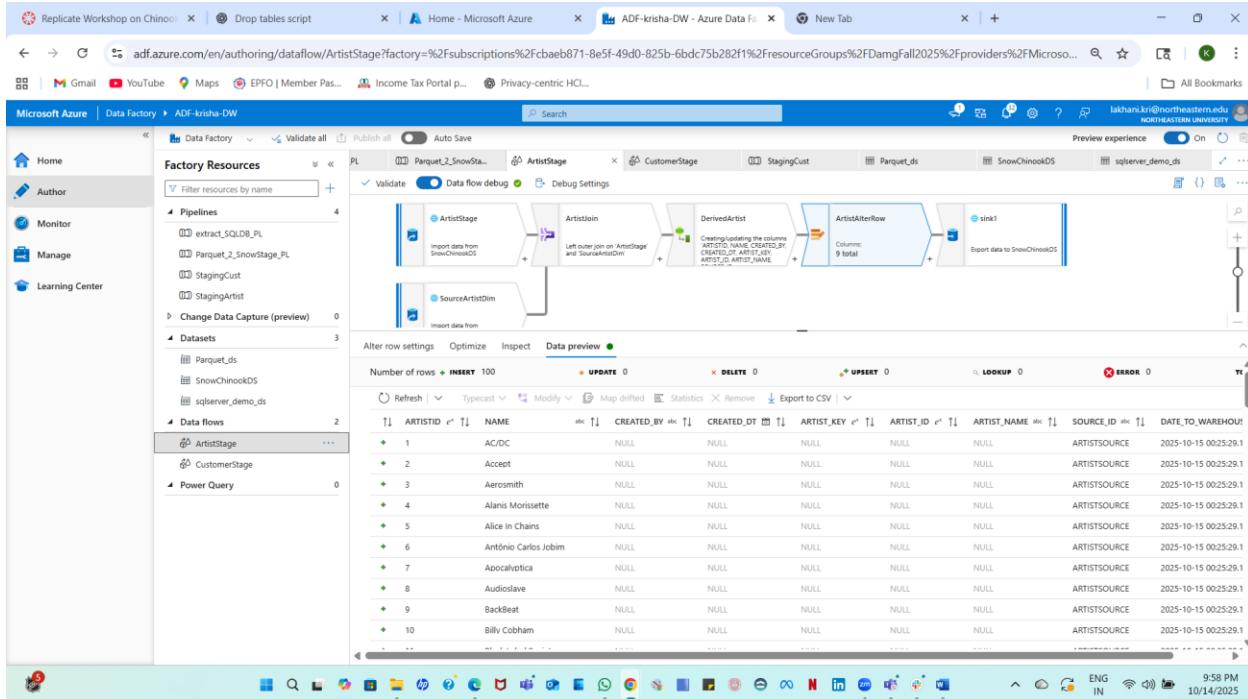
The screenshot shows the 'Data preview' section for the 'ArtistStage' data flow. It displays 100 rows of data from the 'ArtistStage' source. The columns are ARTISTID, NAME, CREATED_BY, CREATED_DT, ARTIST_KEY, ARTIST_ID, ARTIST_NAME, SOURCE_ID, and DATE_TO_WAREHOUSE. The data includes various artists like AC/DC, Accept, Aerosmith, Alanis Morissette, Alice in Chains, Antonio Carlos Jobim, Apocalyptica, Audioslave, Backbeat, and Billy Cobham, along with their respective details and source information.

ARTISTID	NAME	CREATED_BY	CREATED_DT	ARTIST_KEY	ARTIST_ID	ARTIST_NAME	SOURCE_ID	DATE_TO_WAREHOUSE
1	AC/DC	NULL	NULL	NULL	NULL	NULL	ARTISTSOURCE	2025-10-15 00:25:29.1
2	Accept	NULL	NULL	NULL	NULL	NULL	ARTISTSOURCE	2025-10-15 00:25:29.1
3	Aerosmith	NULL	NULL	NULL	NULL	NULL	ARTISTSOURCE	2025-10-15 00:25:29.1
4	Alanis Morissette	NULL	NULL	NULL	NULL	NULL	ARTISTSOURCE	2025-10-15 00:25:29.1
5	Alice in Chains	NULL	NULL	NULL	NULL	NULL	ARTISTSOURCE	2025-10-15 00:25:29.1
6	Antonio Carlos Jobim	NULL	NULL	NULL	NULL	NULL	ARTISTSOURCE	2025-10-15 00:25:29.1
7	Apocalyptica	NULL	NULL	NULL	NULL	NULL	ARTISTSOURCE	2025-10-15 00:25:29.1
8	Audioslave	NULL	NULL	NULL	NULL	NULL	ARTISTSOURCE	2025-10-15 00:25:29.1
9	Backbeat	NULL	NULL	NULL	NULL	NULL	ARTISTSOURCE	2025-10-15 00:25:29.1
10	Billy Cobham	NULL	NULL	NULL	NULL	NULL	ARTISTSOURCE	2025-10-15 00:25:29.1

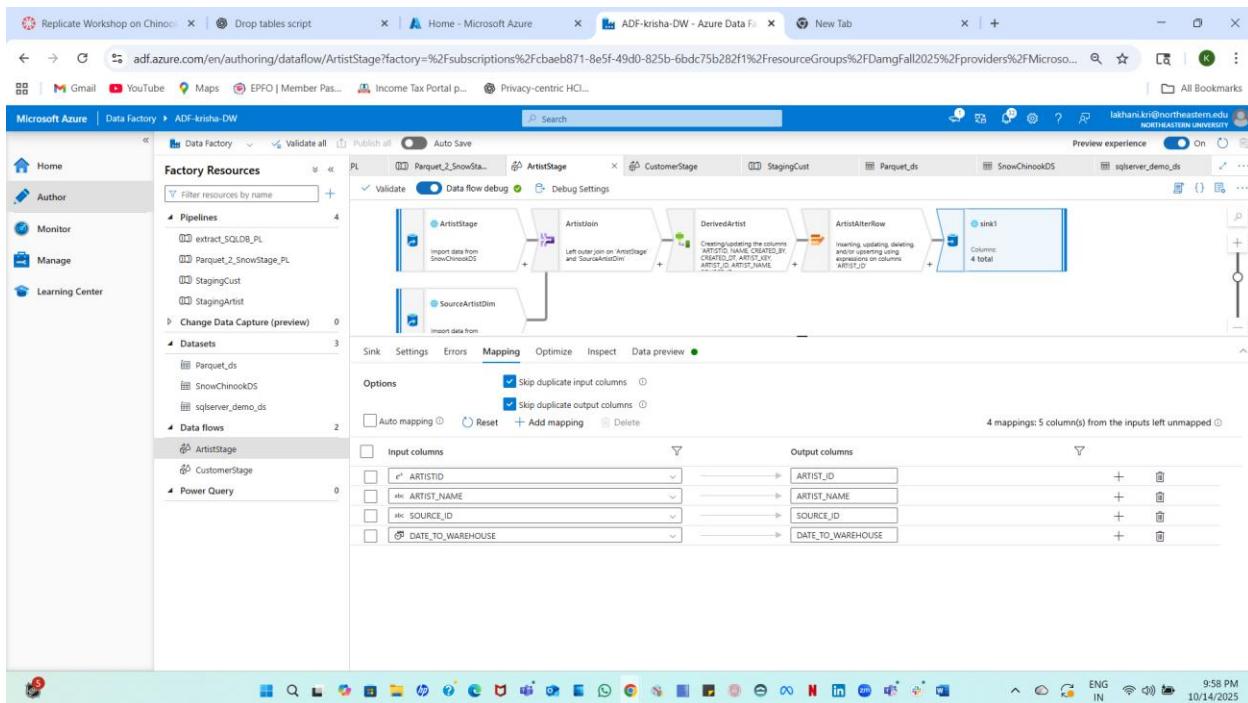
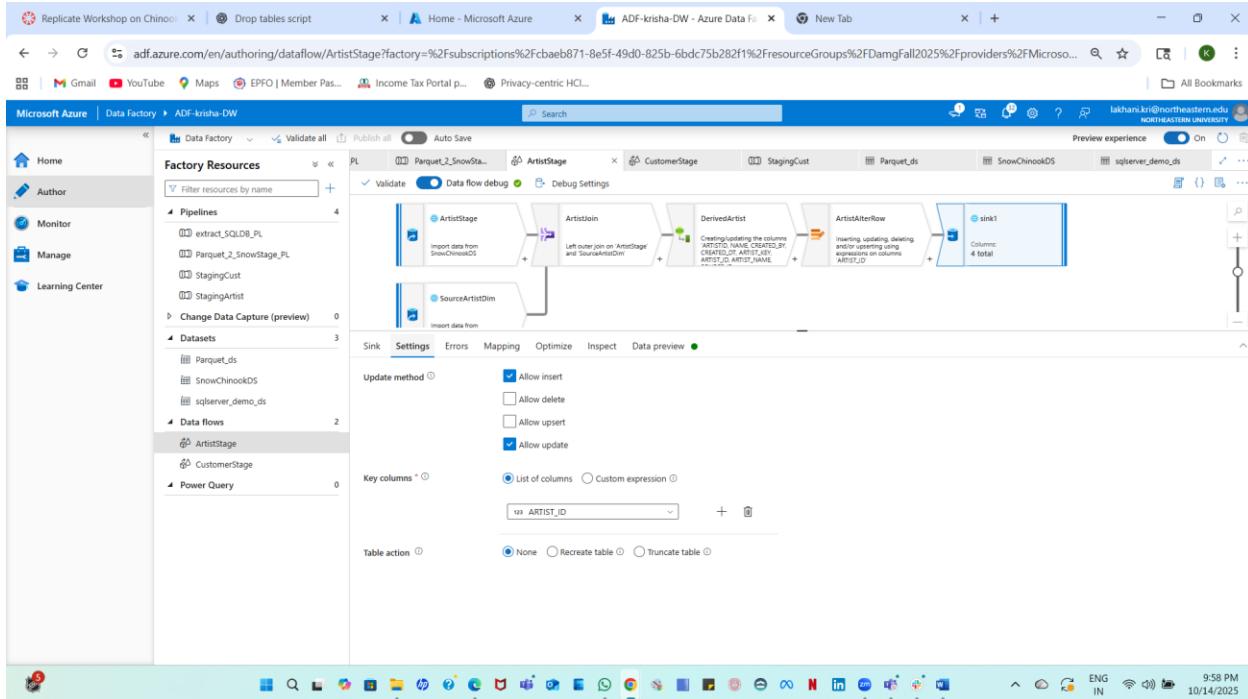
Assignment 5



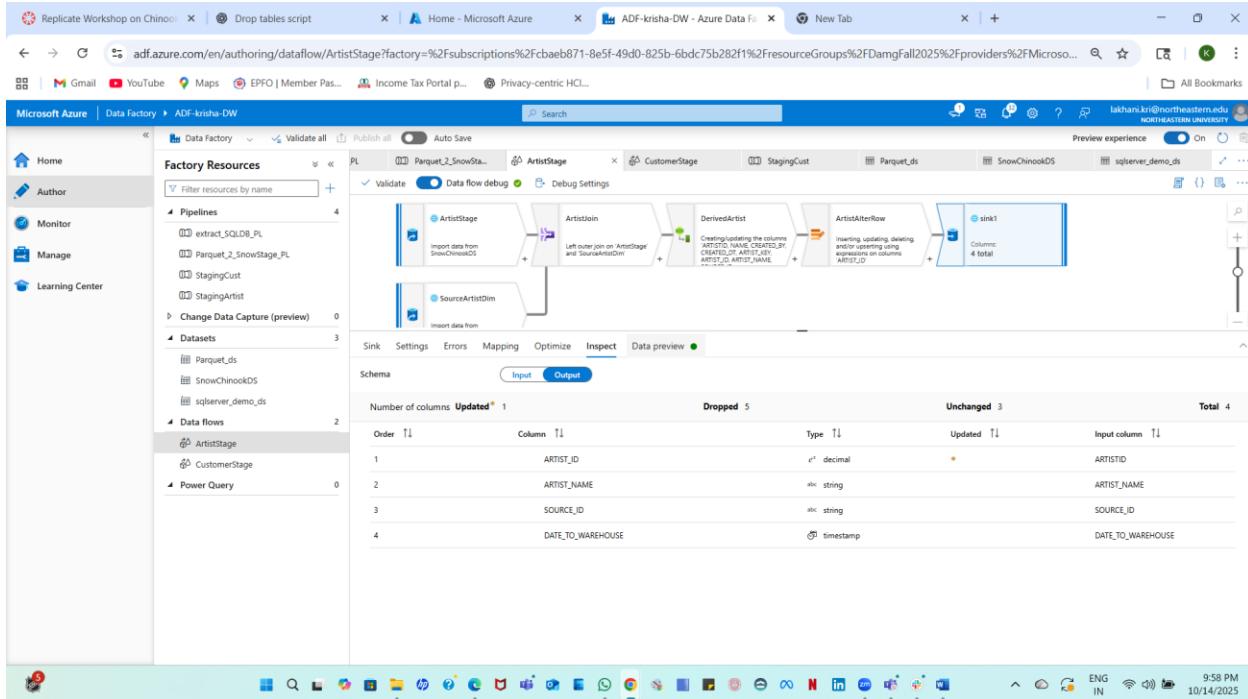
Assignment 5



Assignment 5



Assignment 5



The screenshot shows the Microsoft Azure Data Factory Data Flow interface, identical to the one above, but with the **Data preview** tab selected instead of **Inspect**.

The data preview table shows the results of the data flow execution:

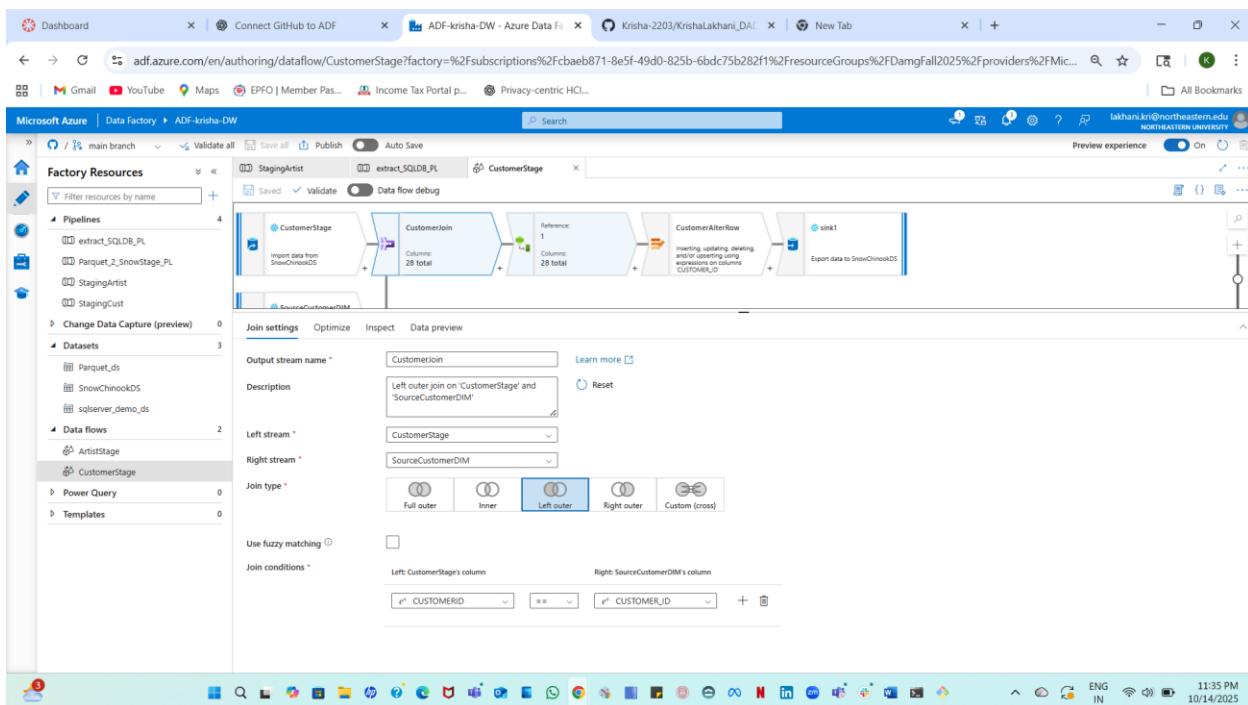
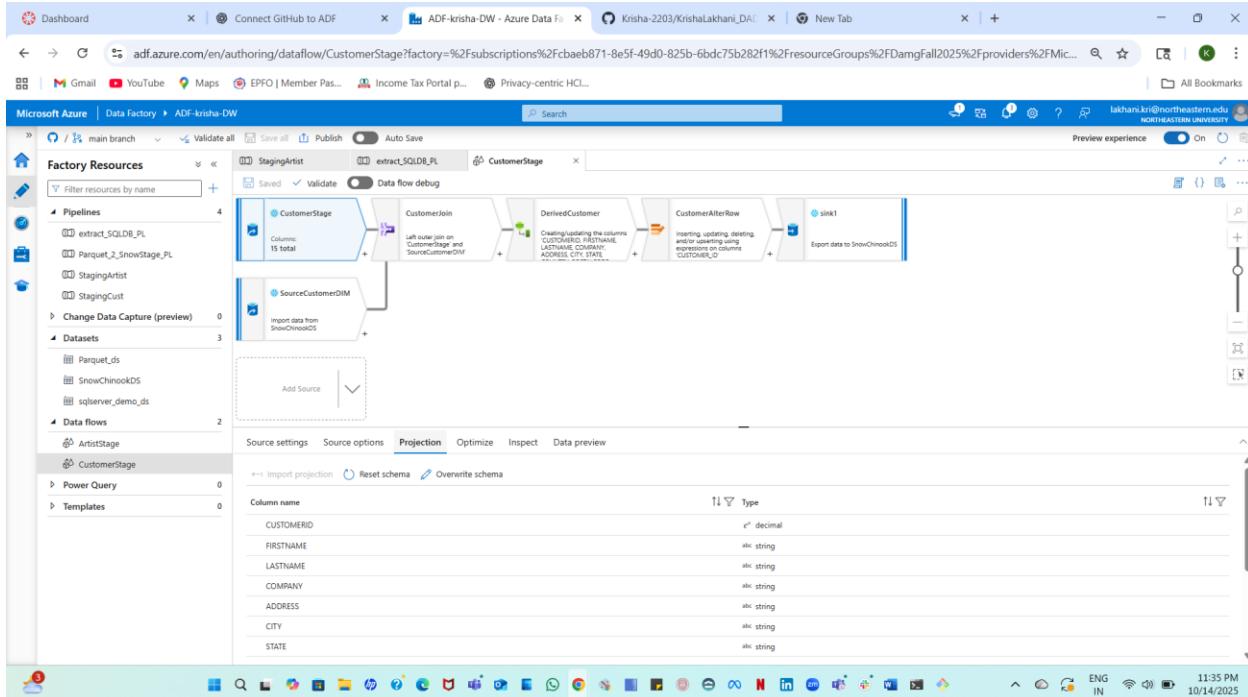
ARTIST_ID	ARTIST_NAME	SOURCE_ID	DATE_TO_WAREHOUSE
1	NULL	ARTISTSOURCE	2025-10-15 00:25:29.195
2	NULL	ARTISTSOURCE	2025-10-15 00:25:29.195
3	NULL	ARTISTSOURCE	2025-10-15 00:25:29.195
4	NULL	ARTISTSOURCE	2025-10-15 00:25:29.195
5	NULL	ARTISTSOURCE	2025-10-15 00:25:29.195
6	NULL	ARTISTSOURCE	2025-10-15 00:25:29.195
7	NULL	ARTISTSOURCE	2025-10-15 00:25:29.195
8	NULL	ARTISTSOURCE	2025-10-15 00:25:29.195
9	NULL	ARTISTSOURCE	2025-10-15 00:25:29.195
10	NULL	ARTISTSOURCE	2025-10-15 00:25:29.195
11	NULL	ARTISTSOURCE	2025-10-15 00:25:29.195

Assignment 5

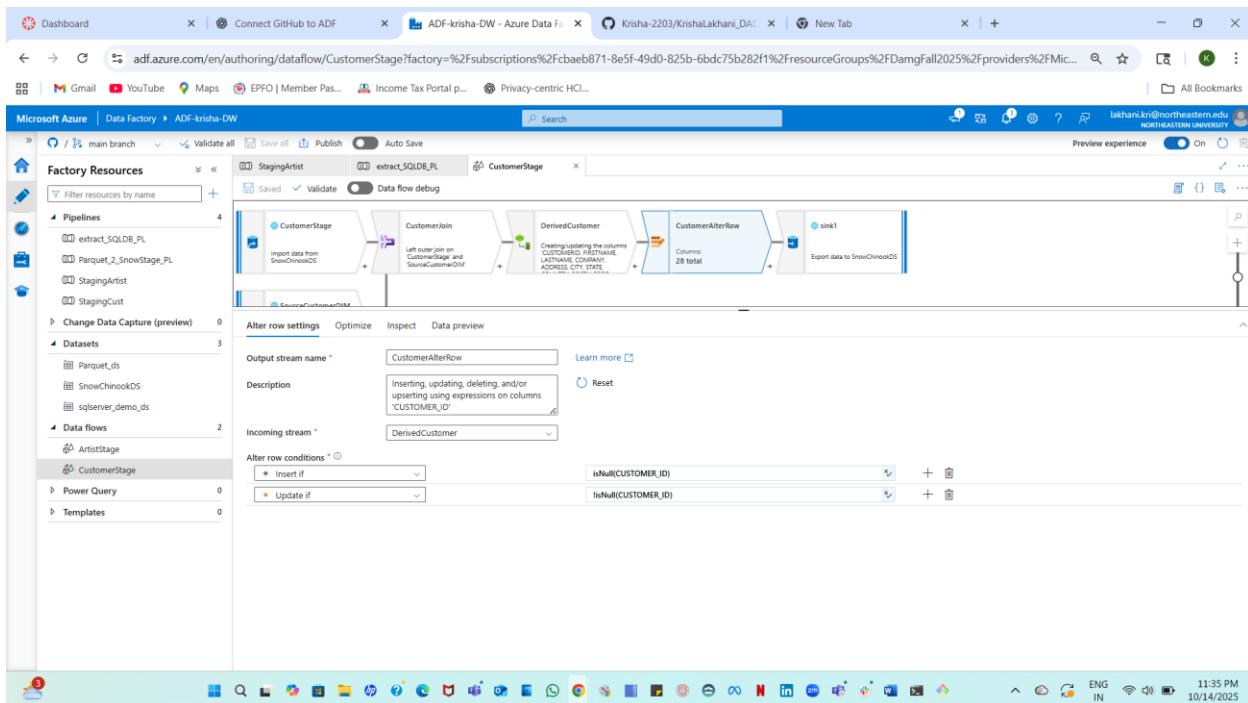
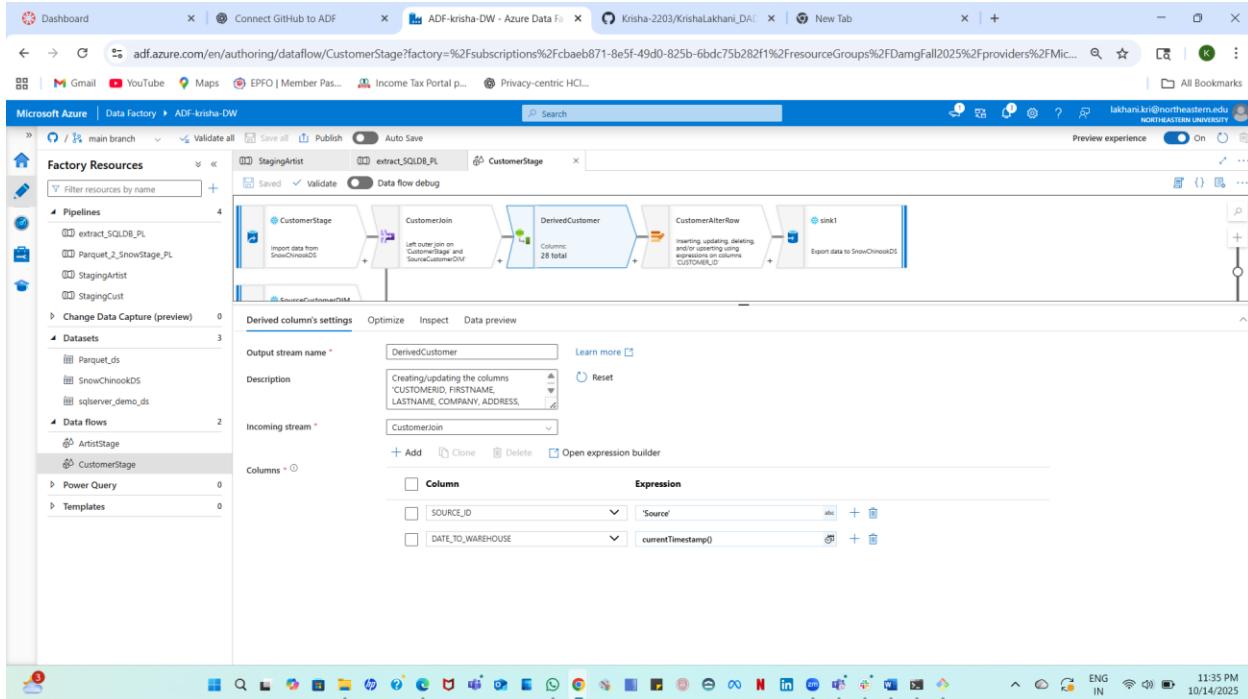
The screenshot shows the Snowflake Data Studio interface. On the left, the Database Explorer displays the schema structure, including the MEDIA_DB database and its tables like ARTIST_DIM, DATE_DIM, and CUSTOMER_DIM. In the center, a query result table for the ARTIST_DIM table is shown with columns ARTIST_KEY, ARTIST_ID, ARTIST_NAME, SOURCE_ID, and DATE_TO_WAREHOUSE. The results show 275 rows with various values for each column. On the right, the Object Details panel shows the definition of the MEDIA_DB.DW.ARTIST_DIM table with columns ARTIST_KEY, ARTIST_ID, ARTIST_NAME, SOURCE_ID, and DATE_TO_WAREHOUSE.

The screenshot shows the Microsoft Azure Data Factory Data Flow blade. It displays a complex data pipeline diagram with multiple stages: CustomerStage, CustomerIn, CustomerAfterRow, and link1. The CustomerStage stage includes a SourceCustomerDIM dataset from SnowChinookDS. The CustomerIn stage performs a Left outer join between CustomerStage and SourceCustomerDIM. The CustomerAfterRow stage performs an Inserting/updating/deleting and/or truncating using expression operation on the CUSTOMER_ID column. The link1 stage exports data to SnowChinookDS. The left sidebar shows the Data Factory navigation menu and the current pipeline named CustomerStage.

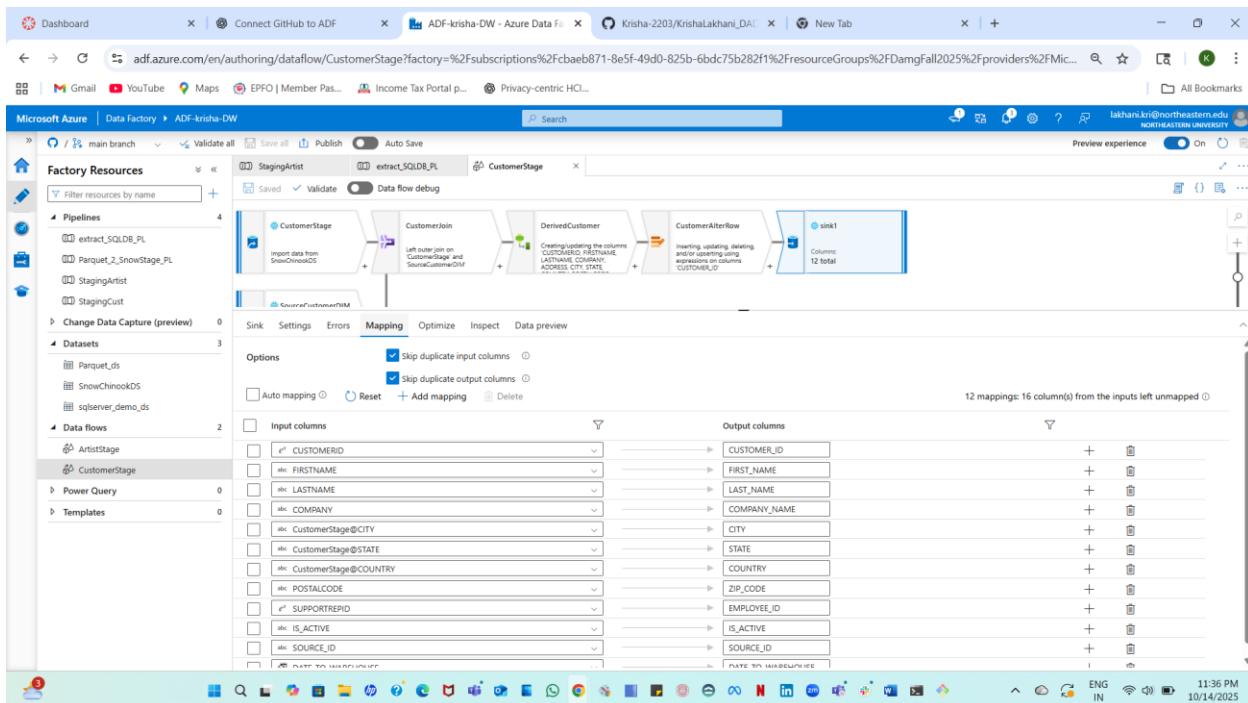
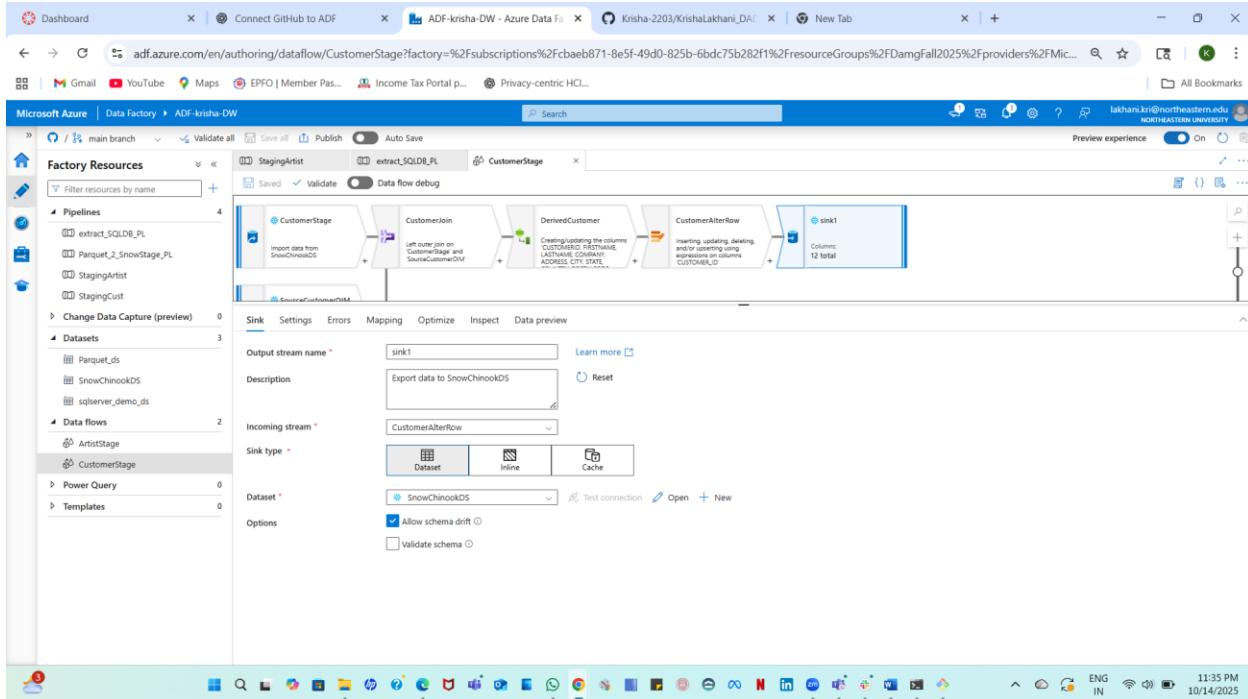
Assignment 5



Assignment 5



Assignment 5



Assignment 5

Screenshot of a Microsoft Edge browser window showing a Snowflake database interface. The URL is <https://app.snowflake.com/east-us-2.azure/dda33793/#/workspaces/ws%24/PUBLIC/DEFAULT%24>.

The left sidebar shows the user's workspace, a search bar, and a Database Explorer pane. The Database Explorer lists the following schema:

- MEDIA_DB**
 - DW**
 - Tables: ARTIST_DIM, CUSTOMER_DIM, DATE_DIM, SALES_FACT, TIME_DIM.
 - Sequences: ARTIST_DIM_SEQ, CUSTOMER_DIM_SEQ, SALES_FACT_SEQ.
 - Information Schema: SAKILA, STAGE.

The main content area displays a query editor with the following code:

```
2
3 | SELECT * from MEDIA_DB.DW.CUSTOMER_DIM;
4 | SELECT * from MEDIA_DB.DW.ARTIST_DIM;
```

The results of the query are shown in a table format:

CUSTOMER_KEY	CUSTOMER_ID	FIRST_NAME	LAST_NAME	COMPANYNAME	CITY	STATE	COUNTRY	ZIP_CODE
1	59	Frank	Almeida	Apple Inc.	—	—	USA	22.05
		Mark	Barnett	Banco do B...	—	SP 5.1%	Can.	001...
				+57 more	+8 more	+51 more	+23 more	+53 more
55	Mark	Taylor	null	Sidney	NSW	Australia	2010	
21	Kathy	Chase	null	Reno	NV	USA	89503	
59	Pita	Srivastava	null	Bangalore	null	India	560001	
16	Frank	Harris	Google Inc.	Mountain Vie	CA	USA	94043-1351	

Below the table, a history of recent queries is listed:

- just now: 76ms SELECT * from MEDIA_DB.DW.CUSTOMER_DIM;
- 1 hour ago: 760ms SELECT * from MEDIA_DB.DW.ARTIST_DIM;
- 1 hour ago: 15s SELECT * from MEDIA_DB.DW.CUSTOMER_DIM;
- 2 hours ago: 25ms SELECT * from MEDIA_DB.DW.CUSTOMER_DIM;
- 2 hours ago: 42ms SELECT * from MEDIA_DB.DW.CUSTOMER_DIM;
- 2 hours ago: 32ms SELECT * from MEDIA_DB.DW.CUSTOMER_DIM;

The bottom right corner shows the system tray with various icons and the date/time: 11/4/2022.

The screenshot shows the Microsoft Azure Data Factory interface for the 'ADF-krisha-DW' dataset. The left sidebar lists resources: Pipelines (extract_SQLDB_PL, Parquet_2_SnowStage_PL, StagingArtist, StagingCust), Datasets (Parquet_Lds, SnowChinookDS, sqsviewer_demo_ds), Data flows (ArtistStage, CustomerStage), and Power Query. The main pane displays the 'extract_SQLDB_PL' pipeline. It shows a 'ForEach' activity with a single item 'ForEachSourceTable'. Below the pipeline details, a table lists four activities with their status, run start time, duration, integration runtime, user properties, and activity run ID.

Activity name	Activity st...	Activit...	Run start	Duration	Integration runtime	User prop...	Activity run ID
sql_2_parquet	Succeeded	Copy data	10/14/2025, 11:37:57 PM	21s	integrationRuntime1		ec798718-f89e-4b2c-81dd-f516caaf
sql_2_parquet	Succeeded	Copy data	10/14/2025, 11:37:33 PM	23s	integrationRuntime1		7385bd957-bd8e-4e24-be09-0166ba
sql_2_parquet	Succeeded	Copy data	10/14/2025, 11:37:33 PM	23s	integrationRuntime1		763flea7-3729-4ec-92c3-6deec5e4f
ForEachSourceTable	Succeeded	ForEach	10/14/2025, 11:37:33 PM	48s			e7d5f90f-c098-40fb-9386-34dc2355

Assignment 5

The screenshot shows the Azure Data Factory interface for the pipeline 'extract_SQLDB_PL'. The left sidebar lists various resources: Pipelines, Datasets, Data flows, and Power Query. The main workspace displays the 'Activities' section for the 'Parquet_2_SnowStage_PL' pipeline. A 'ForEach' activity is selected, which contains a 'Parquet_2_SnowStage' activity. The pipeline status is 'Succeeded'. The run history table shows four runs:

Activity name	Activity st...	Activit...	Run start	Duration	Integration runtime	User prop...	Activity run ID
Parquet_2_Snow_Stage	Succeeded	Copy data	10/14/2025, 11:40:02 PM	38s	AutoResolveIntegrationRuntime (East US 2) AutoR	8b40c54e-58af-4580-a71f-075c304bb5cc	f8b435f47-f5ad-489b-a6f1-f01165a9e
Parquet_2_Snow_Stage	Succeeded	Copy data	10/14/2025, 11:39:17 PM	42s	AutoResolveIntegrationRuntime (East US 2) AutoR	52ff1d15f-de09-4780-aff2-942f1e07f8	c87cb0c1-e71b-40e7-9d55-42fe6350
Parquet_2_Snow_Stage	Succeeded	Copy data	10/14/2025, 11:39:17 PM	43s	AutoResolveIntegrationRuntime (East US 2) AutoR	f5737c5a-cebc-470b-b7ff-13252100f	f5737c5a-cebc-470b-b7ff-13252100f
ForEachTable	Succeeded	ForEach	10/14/2025, 11:39:17 PM	1m 25s			

The screenshot shows the Azure Data Factory interface for the pipeline 'StagingArtist'. The left sidebar lists various resources: Pipelines, Datasets, Data flows, and Power Query. The main workspace displays the 'Activities' section for the 'Parquet_2_SnowStage_PL' pipeline. A 'Data flow' activity is selected, which contains an 'ArtistStage' activity. The pipeline status is 'Succeeded'. The run history table shows one run:

Activity name	Activity st...	Activit...	Run start	Duration	Integration runtime	User prop...	Activity run ID
ArtistStage	Succeeded	Data flow	10/14/2025, 11:41:25 PM	2m 12s	debugpool-8Cores-General-97 (East US 2)		da65212f-ca56-43d2-aa68-00fb1152

Assignment 5

The screenshot shows the Azure Data Factory pipeline editor interface. On the left, the 'Factory Resources' sidebar lists Pipelines, Datasets, and Data flows. In the center, the 'extract_SQLDB_PL' pipeline is selected, showing three stages: 'StagingCust', 'StagingArtist', and 'Parquet_2_SnowStage'. The 'CustomerStage' activity within the 'StagingCust' stage is highlighted. The 'Output' tab is selected, displaying a table of pipeline run details:

Activity name	Activity st...	Activit...	Run start	Duration	Integration runtime	User prop...	Activity run ID
CustomerStage	Succeeded	Data flow	10/14/2025, 11:44:26 PM	1m 26s	debugpool-8Cores-General-97 (East US 2)	d16f3b93-2fb2-4a41-bb84-729f008c	

The status bar at the bottom indicates the system is in English (ENG IN) and the date is 10/14/2025.

Github link - https://github.com/Krishna-2203/KrishnaLakhani_DADABI_Assignments